

# Operating Systems Final Study Guide

## Main Memory

### Memory Management

Memory management involves managing the main memory to allocate and deallocate spaces for processes and data. Key points include:

- **Responsibilities:**
  - Keep track of memory usage (used/free spaces).
  - Allocate memory efficiently to processes.
  - Ensure protection and isolation of processes.
  - Handle deallocation and compaction to reduce fragmentation.


### Allocation Strategies (Assignment 10)

- **First-Fit:** Assign the first hole that fits the process.
- **Best-Fit:** Assign the smallest available hole that fits.
- **Worst-Fit:** Assign the largest available hole.

### Logical vs. Physical Address Space

- **Logical Address:**
  - Generated by the CPU during program execution.
  - Also called virtual address.
- **Physical Address:**
  - Actual location in memory.
  - Managed by the Memory Management Unit (MMU), which translates logical addresses to physical addresses.
- **Address Binding:**
  - Can occur at compile-time, load-time, or run-time.

# Swapping



- **Definition:** A process is temporarily moved out of main memory to secondary storage (disk) to free up space for other processes.
- **Use:** Improves multiprogramming by allowing more processes to run concurrently.
- **Overhead:** Disk I/O operations are slow, so frequent swapping can degrade performance.

Swapping is efficient only for idle or partially idle processes due to the following reasons:

1. **Minimized Interruptions:**

- Swapping out an active process disrupts its execution, leading to significant performance degradation.
- Idle or partially idle processes have minimal CPU or I/O activity, so swapping them out has a lesser impact on the system's responsiveness and throughput.

2. **Avoiding I/O Conflicts:**

- Active processes often rely on continuous I/O operations. Swapping such processes out can lead to **double buffering**, where the operating system must

temporarily buffer I/O data. This adds overhead and slows down both the swapped process and system performance.

3. **Swapping Overhead:**

- The time required to swap a process depends on its memory size and disk transfer rate. For an active process, this overhead delays critical operations.
- Idle processes, on the other hand, do not require immediate CPU or memory access, making the overhead less disruptive.

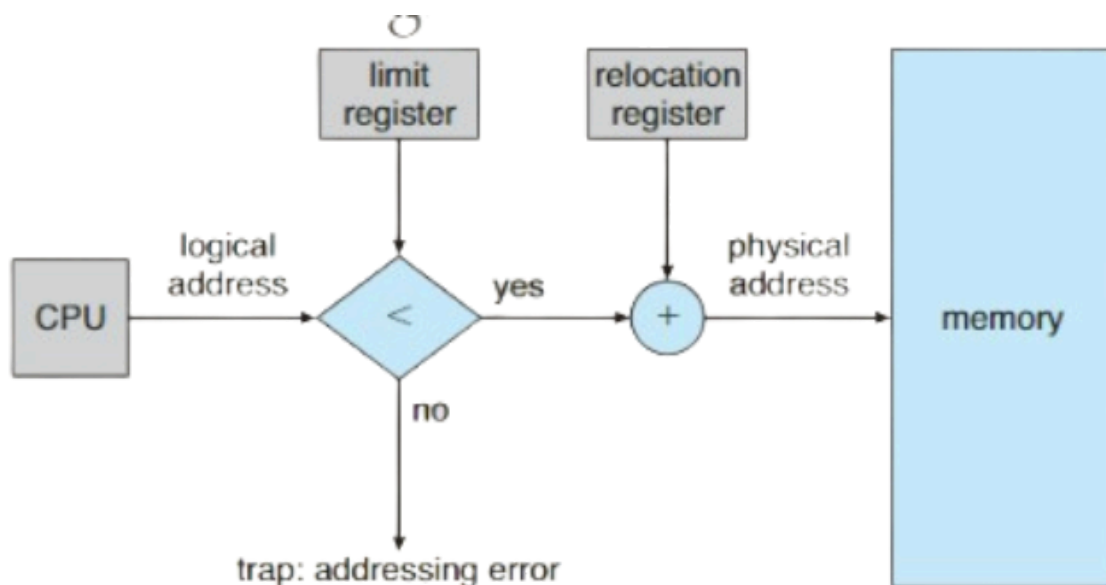
4. **Resource Optimization:**

- Swapping out idle processes frees up memory for active processes, improving resource allocation and utilization.

5. **Avoiding Thrashing:**

- If multiple active processes are swapped in and out frequently, the system may enter a state of **thrashing**, where it spends more time swapping than executing processes. Keeping active processes in memory reduces this risk.

## Memory Protection



- Prevents processes from interfering with each other's memory space.
- **Techniques:**
  - **Base and Limit Registers:** Define the range of accessible addresses.
  - **Segmentation:** Divides memory into segments, each with permissions.
  - **Paging:** Ensures each process can only access its allocated pages.

This diagram explains **Memory Protection** in the context of operating systems using the **limit register** and **relocation register** mechanism. Here's how it works:

---

## Purpose

Memory protection ensures that processes cannot access memory locations outside their allocated range, preventing unintended interference between processes or with the operating system.

---

## Components Explained

### 1. Logical Address:

- The address generated by the CPU during program execution.
- It is a relative address that needs to be translated to a physical address.

### 2. Relocation Register:

- Holds the base address of the process in physical memory.
- Adds this base address to the logical address to calculate the actual **physical address**.

### 3. Limit Register:

- Specifies the size of the allocated memory range for the process.
- Ensures that the process can only access memory within its allocated range.

### 4. Memory Access Check:

- The CPU compares the logical address to the value in the **limit register**.
  - If the logical address is less than or equal to the limit, the address is valid.
  - If the address is outside the limit, the system triggers a **trap** (an error or interrupt) to handle the invalid access.
- 

## Process during Context Switch

### 1. During a **context switch**:

- The operating system (OS) dispatcher sets the **relocation register** and **limit register** values for the next process.
- This ensures that each process operates within its memory bounds and cannot access another process's memory space.

### 2. After the registers are set:

- Any memory access performed by the process will automatically be checked against these registers by the hardware.

---

## Why It's Important

1. **Security:** Prevents processes from accessing memory that they don't own, which could lead to data corruption or breaches.
2. **Isolation:** Ensures that processes are isolated from each other.
3. **Fault Tolerance:** If a process tries to access an invalid address, the system can safely trap the error and handle it without crashing.

Let me know if you need further clarification!

## Fragmentation

- **Definition:** Wasted memory space due to allocation inefficiencies.

### Internal Fragmentation

**What it is:**

- Memory is allocated in fixed-size blocks, but processes often don't use the entire block, leaving some space unused within the allocated memory.

**Example:**

- Fixed block size: 4KB.
- Process request: 3.7KB.
- Wasted memory: 0.3KB.

**Key Issue:**

- Unused memory within allocated blocks adds up across multiple processes.

### External Fragmentation

**What it is:**

- Free memory is available, but it's broken into non-contiguous chunks. As a result, a process requiring a large block of memory can't be allocated even though enough total memory exists.

**Example:**

- Available memory blocks: **100KB, 150KB, 200KB** (non-contiguous).

- Process request: **250KB**.
- Result: The process **cannot be allocated**, even though the total free memory is **450KB**, because no single block is large enough.

**Key Issue:**

- Memory gets wasted because free blocks are scattered.


**50 Percent Rule:**

- For every  $N$  allocated blocks, approximately  $0.5N$  blocks are lost to fragmentation.

## Paging

- Divides memory into fixed-size blocks called **pages**.
- Eliminates external fragmentation but can cause internal fragmentation.

## Page Tables (Assignment 11)



$n=2$  and  $m=4$  32-byte memory and 4-byte pages

- Used to map logical pages to physical frames in memory.
- **Structure:**
  - Each process has its own page table.
  - Contains entries with frame numbers and optional flags (e.g., valid, dirty).

## Addressing for Page Tables

- **Address Translation:**
  - Logical address = Page Number + Offset.
  - Page number indexes the page table to find the frame number.
  - Offset determines the specific address within the frame.
- **Multi-Level Page Tables:**
  - Breaks large page tables into smaller, hierarchical tables.
  - Reduces memory overhead for the page table itself.
- **Translation Lookaside Buffer (TLB):**
  - Cache for page table entries to speed up address translation.

---

# Virtual Memory

Here's a breakdown for **Virtual Memory** concepts with explanations, key points, and potential questions for each section:

---

## Virtual Memory

### Shared Pages

- **What it is:**
  - Multiple processes can share common code or data pages in memory, such as shared libraries.
- **How it works:**
  - Pages are mapped into the virtual address space of each process.
  - Changes in shared pages (if allowed) are reflected across all processes.
- **Benefits:**
  - Reduces memory usage by avoiding duplication of common code.
  - Simplifies inter-process communication by sharing data.

#### Sample Question:

- Explain how shared pages reduce memory usage. Provide an example where two processes share a library.
- 

### Demand Paging

- **What it is:**
  - A lazy loading mechanism where pages are loaded into memory only when referenced, not at process start.
- **How it works:**
  - If a required page is not in memory (a **page fault** occurs), the operating system fetches it from disk.
- **Advantages:**



- Saves memory space by avoiding preloading unused pages.
- **Disadvantages:**
  - Initial page faults can cause performance delays.

#### Sample Question:

- Describe how demand paging works and explain the role of page faults.
- 

## Effective Access Time (EAT)

- **What it is:**
  - The average time required to access memory, considering page faults and TLB (Translation Lookaside Buffer) hits/misses.
- **Formula:** 
$$[\text{EAT} = (1 - p) \times \text{Memory Access Time} + p \times (\text{Page Fault Service Time})]$$
  - $p$  = Probability of a page fault.
  - **Page Fault Service Time** includes disk access and possibly replacing a page.

#### Sample Question:

- Calculate EAT for a system with:
    - Memory access time = 100ns.
    - Page fault rate = 0.01.
    - Page fault service time = 10ms.
- 

## Page Replacement

- **What it is:**
  - When a page fault occurs and memory is full, the system replaces an existing page with the new page.
- **Goal:**
  - Minimize the number of page faults by replacing pages intelligently.

## Algorithms (Assignment 12)

1. **FIFO (First-In-First-Out):**
  - Oldest page is replaced.
2. **Optimal:**

- Replaces the page that will not be used for the longest time (requires future knowledge).
- 3. **LRU (Least Recently Used):**
  - Replaces the page that hasn't been used for the longest time.
- 4. **Second Chance:**
  - Similar to FIFO but gives a "second chance" to pages if they are referenced.

**Sample Question:**

- Given a reference string and a frame size, calculate page faults for FIFO, LRU, and Optimal algorithms.

## Local vs. Global

- **Local Replacement:**
  - Each process has its own frame allocation and replaces pages within its own frames.
- **Global Replacement:**
  - Pages are replaced from a common pool of frames shared by all processes.
- **Comparison:**
  - Global replacement offers flexibility but can cause starvation for some processes.

**Sample Question:**

- Compare local and global page replacement strategies. Which is better in a highly loaded system?
- 

## Thrashing

- **What it is:**
  - When a process spends more time handling page faults than executing, causing a drastic performance drop.
- **Cause:**
  - Occurs when the process's working set (set of pages it actively uses) exceeds the available memory.
- **Solution:**
  - Increase memory, or reduce the degree of multiprogramming.

**Sample Question:**

- Explain what thrashing is and discuss two ways to mitigate it.

---

## Working Set

- **What it is:**
  - The set of pages a process actively uses over a time window.
- **Purpose:**
  - Helps estimate the memory demand of a process.
- **Working Set Model:**
  - Tracks recently used pages and ensures the working set is in memory to avoid thrashing.

### Sample Question:

- Describe the working set model and how it prevents thrashing.
- 

## Kernel Memory

- **What it is:**
  - Memory used by the operating system for managing processes, resources, and I/O.
- **Allocation Strategies:**
  1. **Slab Allocator:**
    - Divides memory into slabs, pre-allocated for specific data structures.
  2. **Buddy System:**
    - Allocates memory in powers of 2 for flexibility.
- **Key Consideration:**
  - Kernel memory must be protected and efficiently managed to ensure system stability.


### Sample Question:

- Compare the slab allocator and buddy system for managing kernel memory.
-

# Mass Storage Structure

---

## Hard Disk Structure



### Components:

- **Platters:** Circular disks that store data.
  - **Tracks:** Concentric circles on the platter surface.
  - **Sectors:** Subdivisions of tracks, the smallest storage unit.
  - **Cylinders:** Vertical alignment of tracks across multiple platters.
  - **Read/Write Head:** Moves across tracks to read/write data.
  - **Spindle:** Rotates the platters.
- 
- **Access Time:**
    - **Seek Time:** Time to move the read/write head to the correct track.
    - **Rotational Latency:** Time waiting for the desired sector to rotate under the head.

- **Transfer Time:** Time to transfer data once the sector is under the head.

**Sample Question:**

- Explain the structure of a hard disk and how access time is calculated.
- 

## Disk Scheduling

Disk scheduling determines the order in which disk I/O requests are serviced to optimize performance.

### Algorithms:

1. **FCFS (First-Come-First-Serve):**

- Services requests in the order they arrive.
- Simple but may lead to long seek times.

2. **SSTF (Shortest Seek Time First):**

- Services the request closest to the current head position.
- Reduces seek time but can lead to starvation for far-off requests.

3. **SCAN (Elevator Algorithm):**

- Moves the head in one direction, servicing requests along the way, then reverses direction.
- Reduces starvation compared to SSTF.

4. **C-SCAN (Circular SCAN):**

- Similar to SCAN but only services requests in one direction, then jumps back to the beginning.
- Provides more uniform wait times.

5. **LOOK and C-LOOK:**

- Variants of SCAN and C-SCAN that only move as far as the last request in each direction, skipping unnecessary movements.

**Sample Question:**

- Given a disk queue with requests at tracks 98, 183, 37, 122, 14, 124, 65, and 67 (head starts at 53):
    - Calculate the total seek time for FCFS, SSTF, and SCAN.
    - Which algorithm performs best?
- 

## Calculating Average I/O Time (Assignment 13)

- **Formula:**

$$\text{Average I/O Time} = \text{Average Seek Time} + \text{Average Latency} + \left( \frac{\text{Data to Transfer}}{\text{Transfer Rate}} \right) + \text{Controller Overhead}$$

- **Components:**

1. **Average Seek Time:** Time to position the head to the desired track.
2. **Average Latency:** Time for the desired sector to rotate under the head.
  - For a disk rotating at  $R$  revolutions per minute (RPM): 
$$\text{Average Latency} = \frac{1}{2} \times \frac{60}{R}$$
3. **Transfer Time:** Depends on the transfer rate and the size of the data.
  - $$\text{Transfer Time} = \frac{\text{Data Size}}{\text{Transfer Rate}}$$
4. **Controller Overhead:** Additional time introduced by the disk controller.

### Example Calculation:

- Disk: 7,200 RPM, 5ms seek time, 1Gbps transfer rate, 4KB block size, 0.1ms controller overhead.
  - Average Latency: 
$$\frac{1}{2} \times \frac{60}{7,200} = 4.17 \text{ms}$$
  - Transfer Time: 
$$\frac{4 \times 8 \text{Kb}}{1,000 \text{Mbps}} = 0.032 \text{ms}$$
  - Total I/O Time: 
$$5 + 4.17 + 0.032 + 0.1 = 9.302 \text{ms}$$

### Sample Question:

- Calculate the average I/O time for a disk with:
    - 15,000 RPM
    - Seek time: 3ms
    - Transfer rate: 500MBps
    - Data block size: 64KB
    - Controller overhead: 0.2ms
-