

CS 452 Operating System Concepts

CPU Scheduling

Dr. Denton Bobeldyk

Scheduling Algorithms

- ❖ First-Come First Served
- ❖ Shortest Job First
- ❖ Priority Scheduling
- ❖ Round Robin Scheduling
- ❖ Multilevel Queue Scheduling
- ❖ Multilevel Feedback Queue Scheduling

Scheduling Algorithms

- ❖ **First-Come First Served**
- ❖ Shortest Job First
- ❖ Priority Scheduling
- ❖ Round Robin Scheduling
- ❖ Multilevel Queue Scheduling
- ❖ Multilevel Feedback Queue Scheduling

First Come First Serve

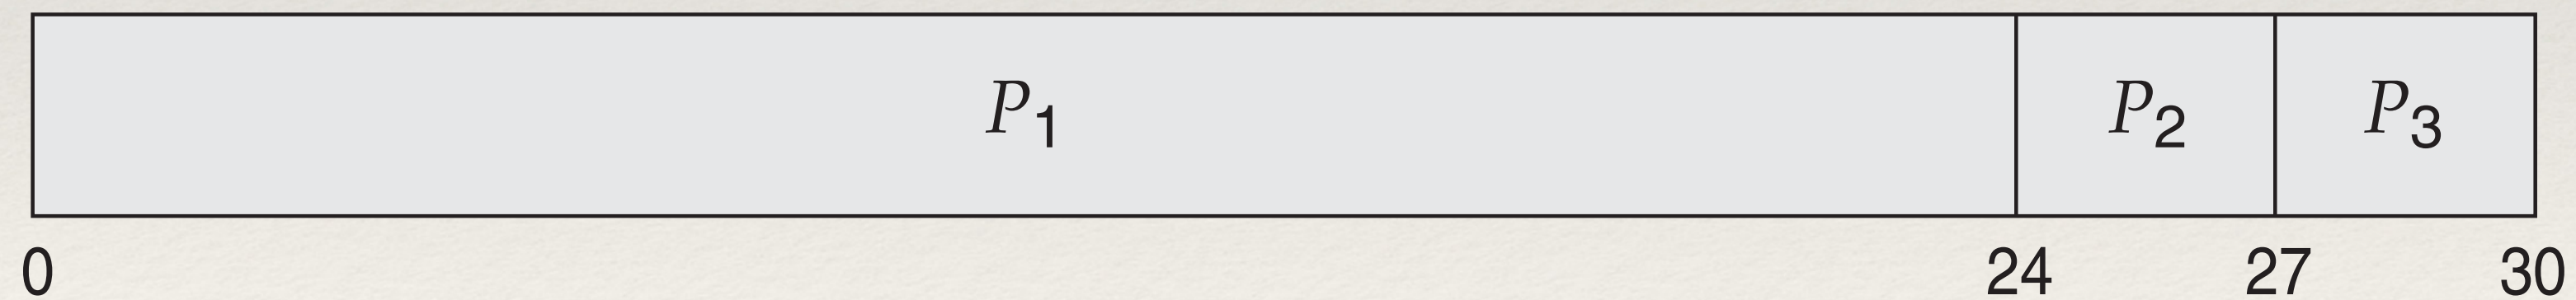
- ❖ First process to request the CPU, gets it 😊
- ❖ Managed with a FIFO queue (First In First Out)

Evaluating FCFS

Process	Burst Time
P1	24
P2	3
P3	3

Evaluating FCFS

Process	Burst Time
P1	24
P2	3
P3	3



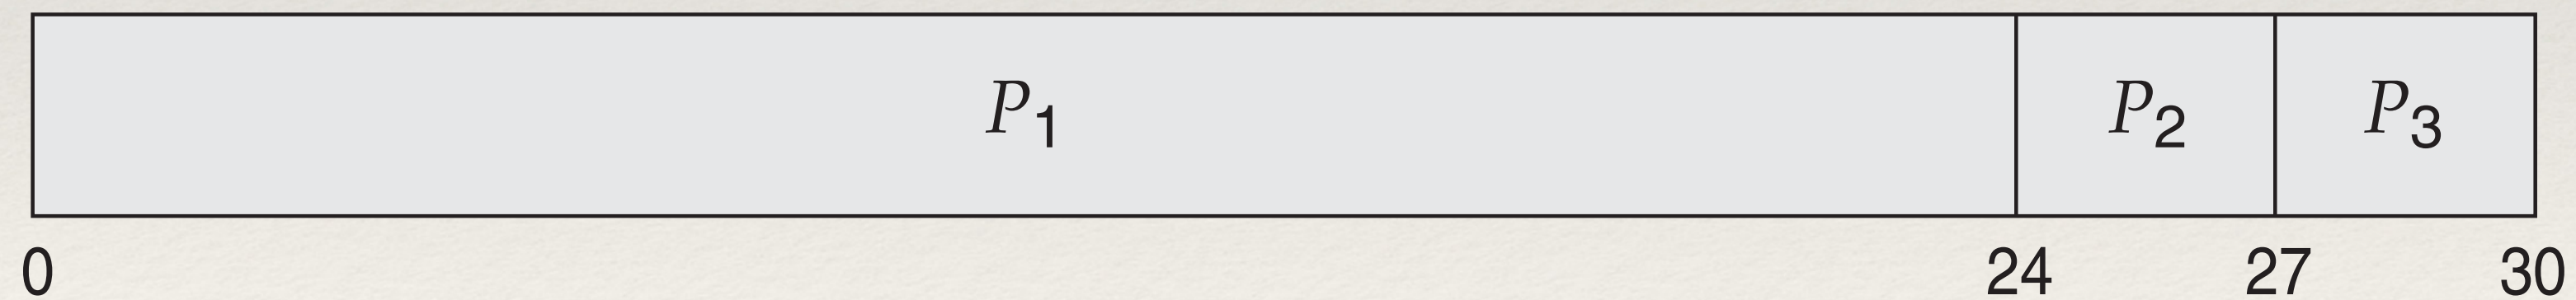
Use of Gant Chart to help visualize!

Scheduling Algorithm Metric

- ❖ Need a metric to compare CPU scheduling algorithms
 - ❖ Calculate the average wait time for each process

Evaluating FCFS

Process	Burst Time
P1	24
P2	3
P3	3




$$\text{Average Wait Time} = (0 + 24 + 27) / 3 = 17 \text{ milliseconds}$$

Evaluating FCFS

What if the processes arrive in a different order?

Process	Burst Time
P2	3
P3	3
P1	24

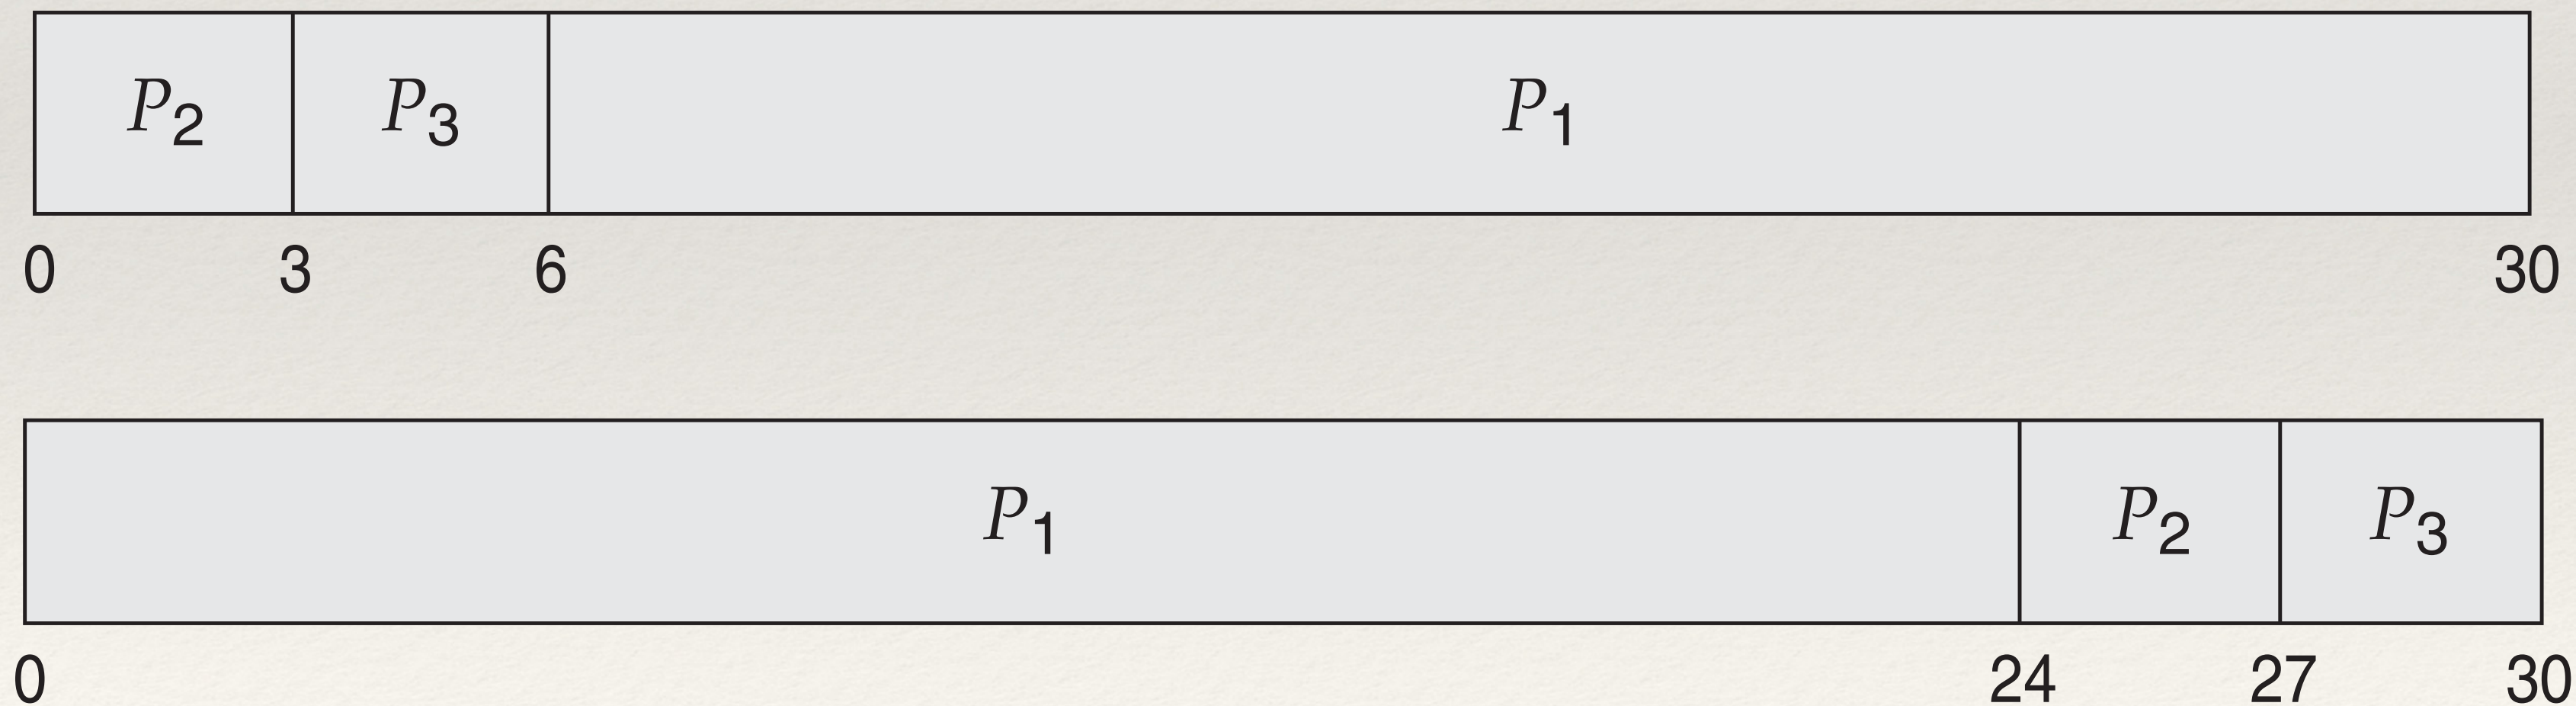


$$\text{Average Wait Time} = (0 + 3 + 6) / 3 = 3 \text{ milliseconds}$$

FCFS Thoughts

FCFS average wait time is heavily dependent on:

1. The burst time of the processes arriving
2. The order in which they arrive



Evaluating FCFS - Individual Exercise

Process	Burst Time
P1	12
P2	1
P3	4

What is the average wait time?

Evaluating FCFS - Individual Exercise

Process	Burst Time
P1	12
P2	1
P3	4

Average Wait Time = $(0 + 12 + 13) / 3 = 8.33$ milliseconds

FCFS Scheduling Algorithm

1. Non preemptive

2. Convoy Effect

- All processes need to wait for one “big” process to complete

Scheduling Algorithms


- ❖ First-Come First Served
- ❖ **Shortest Job First**
- ❖ Priority Scheduling
- ❖ Round Robin Scheduling
- ❖ Multilevel Queue Scheduling
- ❖ Multilevel Feedback Queue Scheduling

Shortest Job First (SJF)

- ❖ When the CPU is available, it is assigned the processes that has the smallest next CPU burst


Shortest Job First (SJF)

Process	Burst Time
P1	6
P2	8
P3	7
P4	3



Shortest Job First (SJF)

Process	Burst Time
P1	6
P2	8
P3	7
P4	3



$$\text{Average Wait Time} = (0 + 3 + 9 + 16) / 4 = 7 \text{ milliseconds}$$

SJF vs FCFS

Process	Burst Time
P1	6
P2	8
P3	7
P4	3

SJF: Average Wait Time = $(0 + 3 + 9 + 16) / 4 = 7$ milliseconds

FCFS: Average Wait Time = $(0 + 6 + 14 + 21) / 4 = 10.25$ milliseconds

SJF Complications

- ❖ What is the burst time for each of the processes arriving at the CPU???

SJF Complications – Group Exercise

- ❖ What is the burst time for each of the processes arriving at the CPU???

How can the burst time be determined?

Calculating Burst Time

- ❖ Predicted as an **exponential average** of the measured length of previous CPU bursts
 - ❖ Let τ_{n+1} be the predicted value of the next burst
 - ❖ Let t_n = length of nth CPU burst
 - ❖ Then for α , $0 \leq \alpha \leq 1$
 - ❖ α determines the weight of recent and past history cpu burst times
 - ❖ $\tau_{n+1} = \alpha t_n + (1 - \alpha)\tau_n$

Calculating Burst Time

- ❖ $\tau_{n+1} = \alpha t_n + (1 - \alpha)\tau_n$
 - t_n contains the most recent information
 - τ_n stores the past history
- ❖ α controls the relative weight of recent and past history
- ❖ If $\alpha = 0$, then recent history has no effect
- ❖ If $\alpha = 1$, then only the most recent CPU burst is relevant
- ❖ More commonly $\alpha = 1/2$
- ❖ Initial τ (τ_0) can be defined as a constant or as an overall system average

Calculating Burst Time

$$\alpha = 0$$

$$\tau_{n+1} = 0 \ t_n + (1 - 0)\tau_n$$

$$\tau_{n+1} = \tau_n$$

Calculating Burst Time

$$\alpha = 1$$

$$\tau_{n+1} = 1 \ t_n + (1 - 1)\tau_n$$

$$\tau_{n+1} = t_n$$

Calculating Burst Time


$$\tau_{n+1} = \alpha t_n + (1 - \alpha)\alpha t_{n-1} + \dots + (1 - \alpha)^j \alpha t_{n-j} + \dots + (1 - \alpha)^{n+1} \tau_0$$

t_n contains the most recent information

τ_n stores the past history

Calculating Burst Time


$$\tau_0 = 10, \alpha = 1/2$$



CPU burst (t_i)		6	4	6	4	13	13	13	...
"guess" (τ_i)	10	8	6	6	5	9	11	12	...

Shortest Job First (SJF)

Process	Arrival Time	Burst Time
P1	0	8
P2	1	4
P3	2	9
P4	3	5



Average Wait Time =

SJF with Preemption

Wait Time Per Job = Start Time - Arrival Time

Wait Time for a preempted job:


(Time of first execution - arrival time) +

$$\sum_{2}^n (\text{startTime}_i - \text{endTime}_{i-1})$$

Where n is the number of times the process was preempted + 1

Shortest Job First (SJF)

Process	Arrival Time	Burst Time	Wait Time
P1	0	8	$(0-0) + (10-1) = 9$
P2	1	4	$1-1 = 0$
P3	2	9	$17-2 = 15$
P4	3	5	$5-3 = 2$



Average Wait Time = $9 + 0 + 15 + 2 = 26 / 4 = 6.5$ milliseconds

SJF Algorithm

- ❖ Preemptive or nonpreemptive
- ❖ Preemptive SJF scheduling is sometimes called shortest-remaining-time-first

Scheduling Algorithms

- ❖ First-Come First Served
- ❖ Shortest Job First
- ❖ **Priority Scheduling**
- ❖ Round Robin Scheduling
- ❖ Multilevel Queue Scheduling
- ❖ Multilevel Feedback Queue Scheduling

Priority Scheduling

- ❖ A priority is associated with each process
- ❖ The CPU is allocated to the process with the highest priority

Priority Scheduling - Number Representation

Which do we use?

- ❖ The lower the number, the higher the priority.
- ❖ The higher the number, the higher the priority.

Priority Scheduling - Number Representation


- ❖ **The lower the number, the higher the priority. (Textbook uses this)**
- ❖ The higher the number, the lower the priority.

Priority Scheduling

- ❖ SJF is an example of priority scheduling
 - ❖ The larger the CPU burst size, the lower the priority
 - ❖ The smaller the CPU burst size, the higher the priority

Priority Scheduling

Process	Arrival Time	Burst Time	Priority	Wait Time
P1	0	10	3	6
P2	0	1	1	0
P3	0	2	4	16
P4	0	1	5	18
P5	0	5	2	1



$$\text{Average Wait Time} = 6 + 0 + 16 + 18 + 1 = 41 / 5 = 8.2 \text{ milliseconds}$$

Priority Scheduling - Problem

- ❖ Indefinite Blocking - Starvation can occur

Priority Scheduling - Problem

- ❖ Indefinite Blocking - Starvation can occur
- ❖ How can this be solved?

Priority Scheduling - Problem

- ❖ Indefinite Blocking - Starvation
 - ❖ Can add an 'age' to each process
 - ❖ The longer a process waits, the better their priority becomes

Scheduling Algorithms


- ❖ First-Come First Served
- ❖ Shortest Job First
- ❖ Priority Scheduling
- ❖ **Round Robin Scheduling**
- ❖ Multilevel Queue Scheduling
- ❖ Multilevel Feedback Queue Scheduling

Round Robin Scheduling

- ❖ Ready Queue is FIFO
- ❖ Time Slice (Quantum) is defined
 - ❖ Generally from 10-100ms in length
- ❖ CPU scheduler allocates CPU to each process for one time slice
- ❖ New processes are added to the tail of the queue


Round Robin Scheduling

Process	Arrival Time	Burst Time
P1	0	24
P2	0	3
P3	0	3




Round Robin Scheduling

Process	Arrival Time	Burst Time
P1	0	24
P2	0	3
P3	0	3



$$\text{Average Wait Time} = (10-4) + 4 + 7 = 17/3 = 5.66 \text{ milliseconds}$$

Round Robin Scheduling



Round Robin Scheduling

- ❖ If the time quantum is really high - same as FCFS
- ❖ If the time quantum is really small - a lot of overhead doing process switching
- ❖ Rule of Thumb: 80 percent of the CPU bursts should be less than the time slice / quantum (no context switching for those processes)

Scheduling Algorithms

- ❖ First-Come First Served
- ❖ Shortest Job First
- ❖ Priority Scheduling
- ❖ Round Robin Scheduling
- ❖ **Multilevel Queue Scheduling**
- ❖ Multilevel Feedback Queue Scheduling

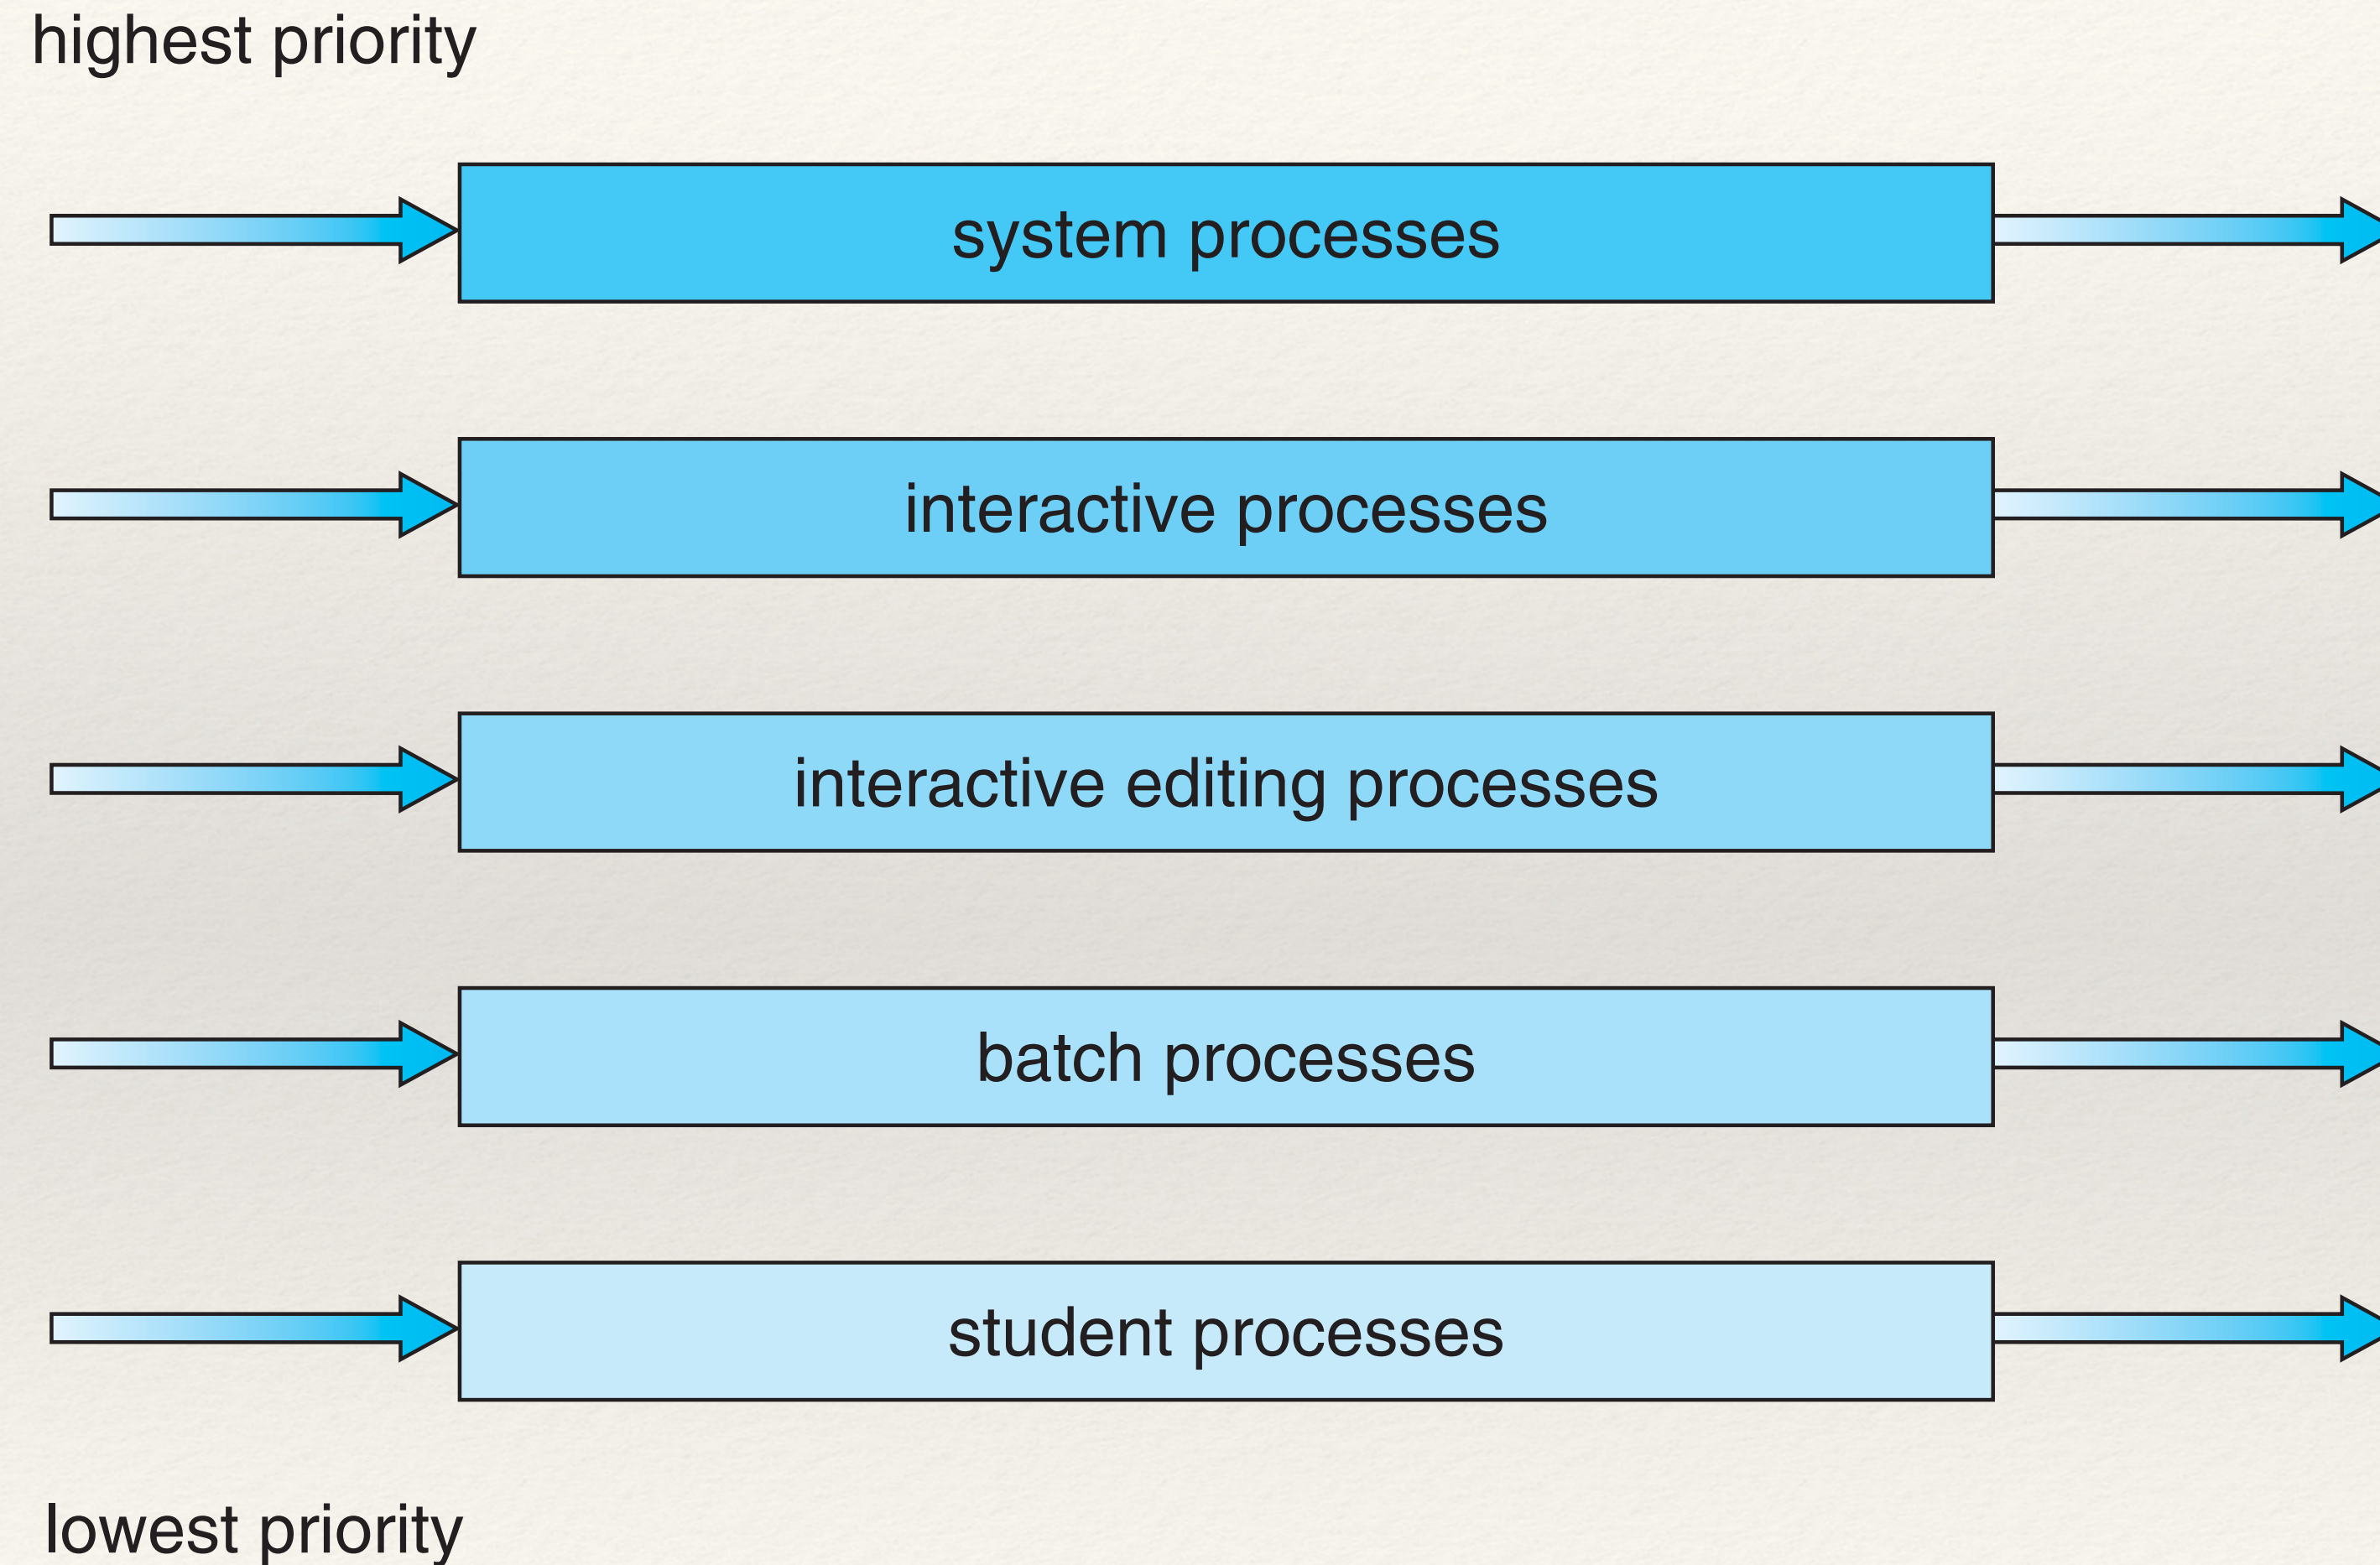
Multi-level Queue Scheduling

- ❖ Ready Queue is partitioned into separate queues:
 - ❖ Foreground (interactive)
 - ❖ Background (batch)
- ❖ Each Queue has its own scheduling algorithm
 - ❖ Foreground - round robin
 - ❖ Background - FCFS

Multi-level Queue Scheduling

- ❖ CPU Scheduling must occur between the queues
 - ❖ Fixed priority scheduling
 - ❖ Serve all foreground, then background
 - ❖ Time slice
 - ❖ 80% to foreground using RR
 - ❖ 20% to background using FCFS

Multilevel queue scheduling



Scheduling Algorithms

- ❖ First-Come First Served
- ❖ Shortest Job First
- ❖ Priority Scheduling
- ❖ Round Robin Scheduling
- ❖ Multilevel Queue Scheduling
- ❖ **Multilevel Feedback Queue Scheduling**

Multi-level Feedback Scheduling

- ❖ Allows processes to move between queues
- ❖ Separate processes by their CPU burst size
 - ❖ If it uses too much time, move to a lower priority queue

Multi-level Feedback Scheduling

- ❖ Schedule all processes in queue 0
- ❖ Once queue 0 is empty, process those in queue 1
- ❖ Once queue 0 and 1 are empty, process those in queue 2

Multi-level Feedback Scheduling

- ❖ Defined by:
 - ❖ The number of queues
 - ❖ The scheduling algorithm for each queue
 - ❖ Method used to determine initial queue
 - ❖ Method used to determine promotion / demotion to a different queue

Multiple-Processor Scheduling

- ❖ Multiple CPUs available
- ❖ Load share over each of the CPUs
- ❖ Scheduling can be more complex
- ❖ Define **homogenous** as processors that are identical in terms of their functionality

Multi-Processor Scheduling

- ❖ Asymmetric multiprocessing
- ❖ Symmetric multiprocessing

Multi-Processor Scheduling

- ❖ **Asymmetric multiprocessing**
 - ❖ Master server (one processor)
 - ❖ Scheduling decisions, I/O processing, System activities
 - ❖ Other processors execute user only code
 - ❖ Reduces need for sharing data across processors
- ❖ **Symmetric multiprocessing**

Multi-Processor Scheduling

- ❖ Asymmetric multiprocessing
- ❖ Symmetric multiprocessing
 - ❖ Each processor examines the ready queue and selects a process to execute
 - ❖ Windows / Linux / Mac OS X use this type of scheduling

Symmetric multiprocessing (SMP)

- ❖ Processor Affinity
- ❖ Load Balancing


Processor Affinity

- ❖ Each processor can have cache memory
- ❖ If a process runs on a processor, the cache memory will be populated with data most recently accessed for that process
- ❖ If the process switches to another processor, the data cache from the original processor must be invalidated and any new data must be loaded into the cache for the new processor
- ❖ It is advantageous for a process to stay on the same processor, aka Processor Affinity

Processor Affinity

- ❖ Soft Affinity - attempt to keep on same processor, but no guarantee
- ❖ Hard Affinity - a process can specify a subset of processors on which it may run
- ❖ Linux provides the 'sched_setaffinity()' system call which supports hard affinity. Default is soft affinity

Processor Affinity



A visual of a CPU accessing slower access memory

Load Balancing

- ❖ Attempt to keep workload even across all processors
- ❖ Push Migration - A specific task periodically checks the load on each processor and evenly distributes the load by moving (or pushing) processes from overloaded processors to idle or less busy processors
- ❖ Pull Migration - An idle processor pulls a waiting task from a busy processor

Load Balancing

- ❖ Linux scheduler implements both techniques
 - ❖ Push & Pull

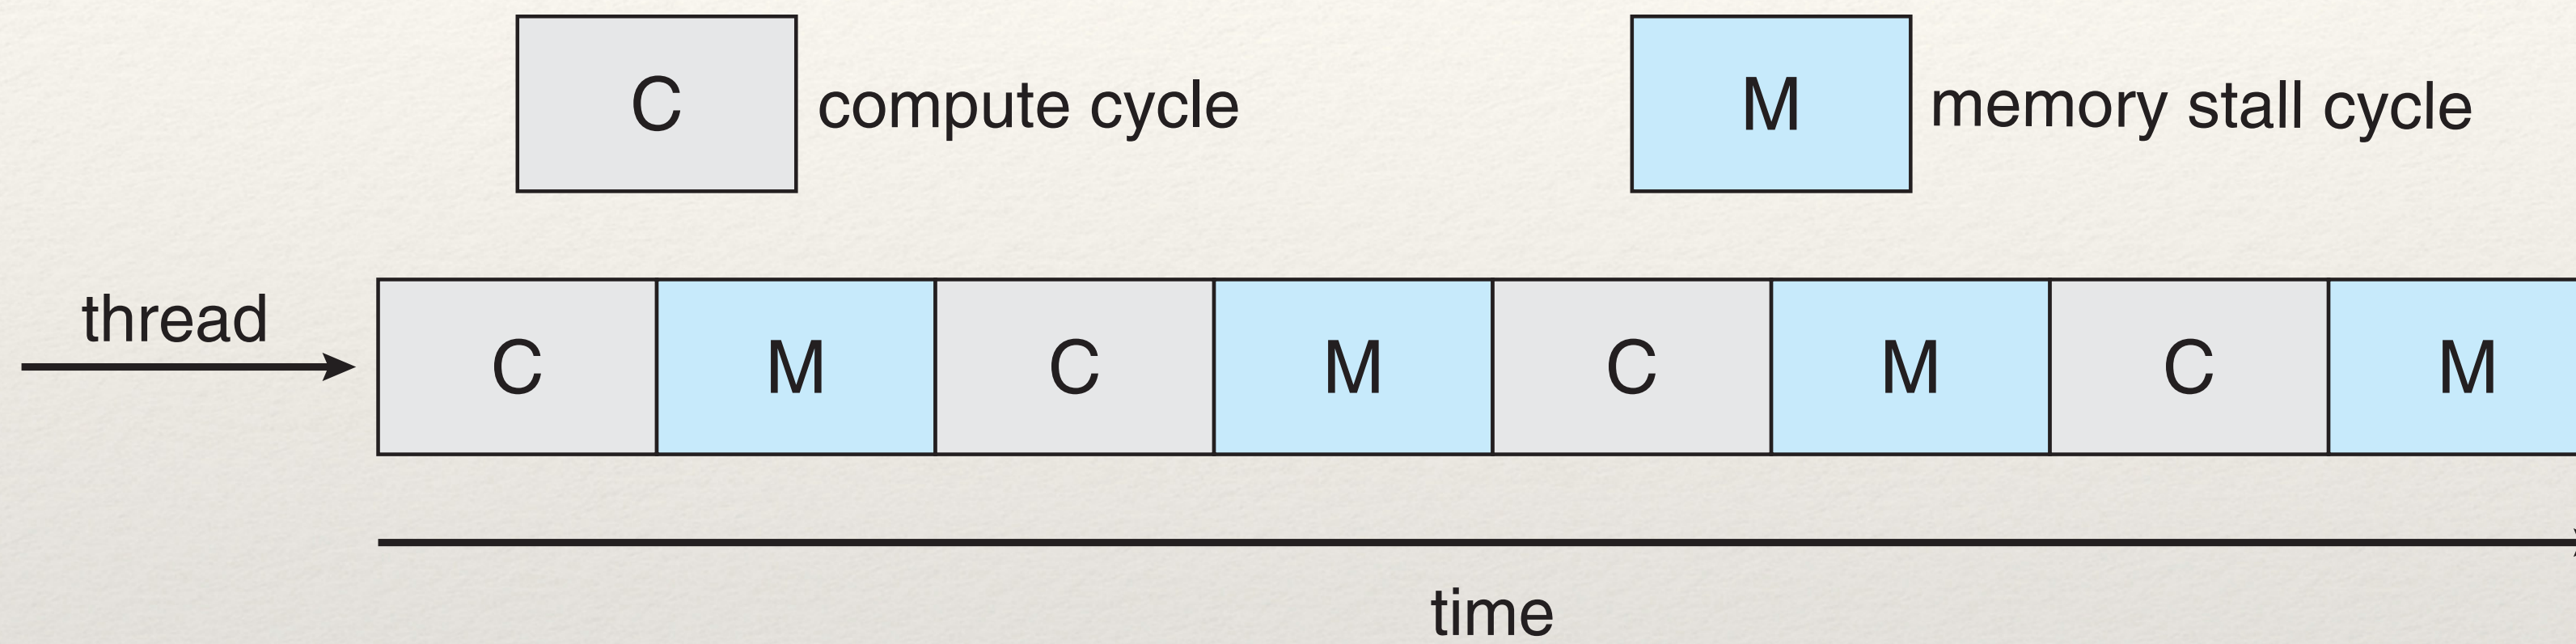
Multicore Processor

- ❖ Multiple processor cores on the same chip
- ❖ OS treats them as individual processor cores
- ❖ Faster and consume less power
- ❖ Complicates scheduling issues

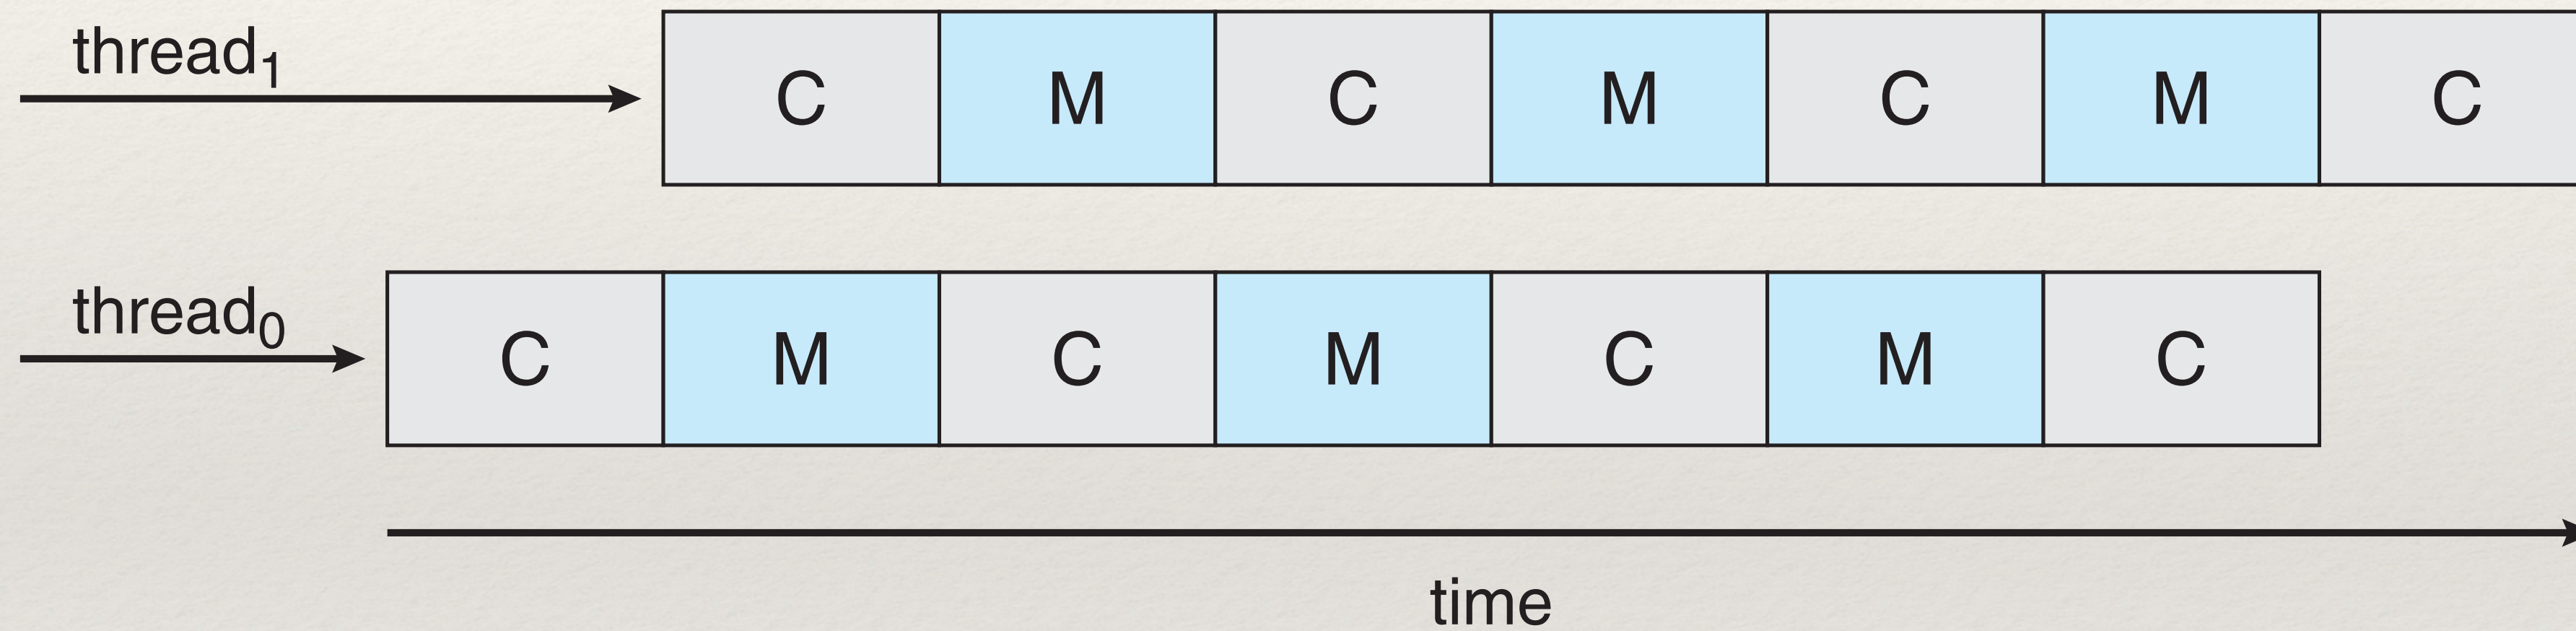
Multicore Processor

- ❖ Processors may spend a significant amount of time waiting for data to become available, this is called **memory stall**
- ❖ Multithreaded processor core: 2 or more hardware threads are assigned to a core
- ❖ Operating system treats each hardware thread as a logical processor

Multicore Processor



Multicore Processor

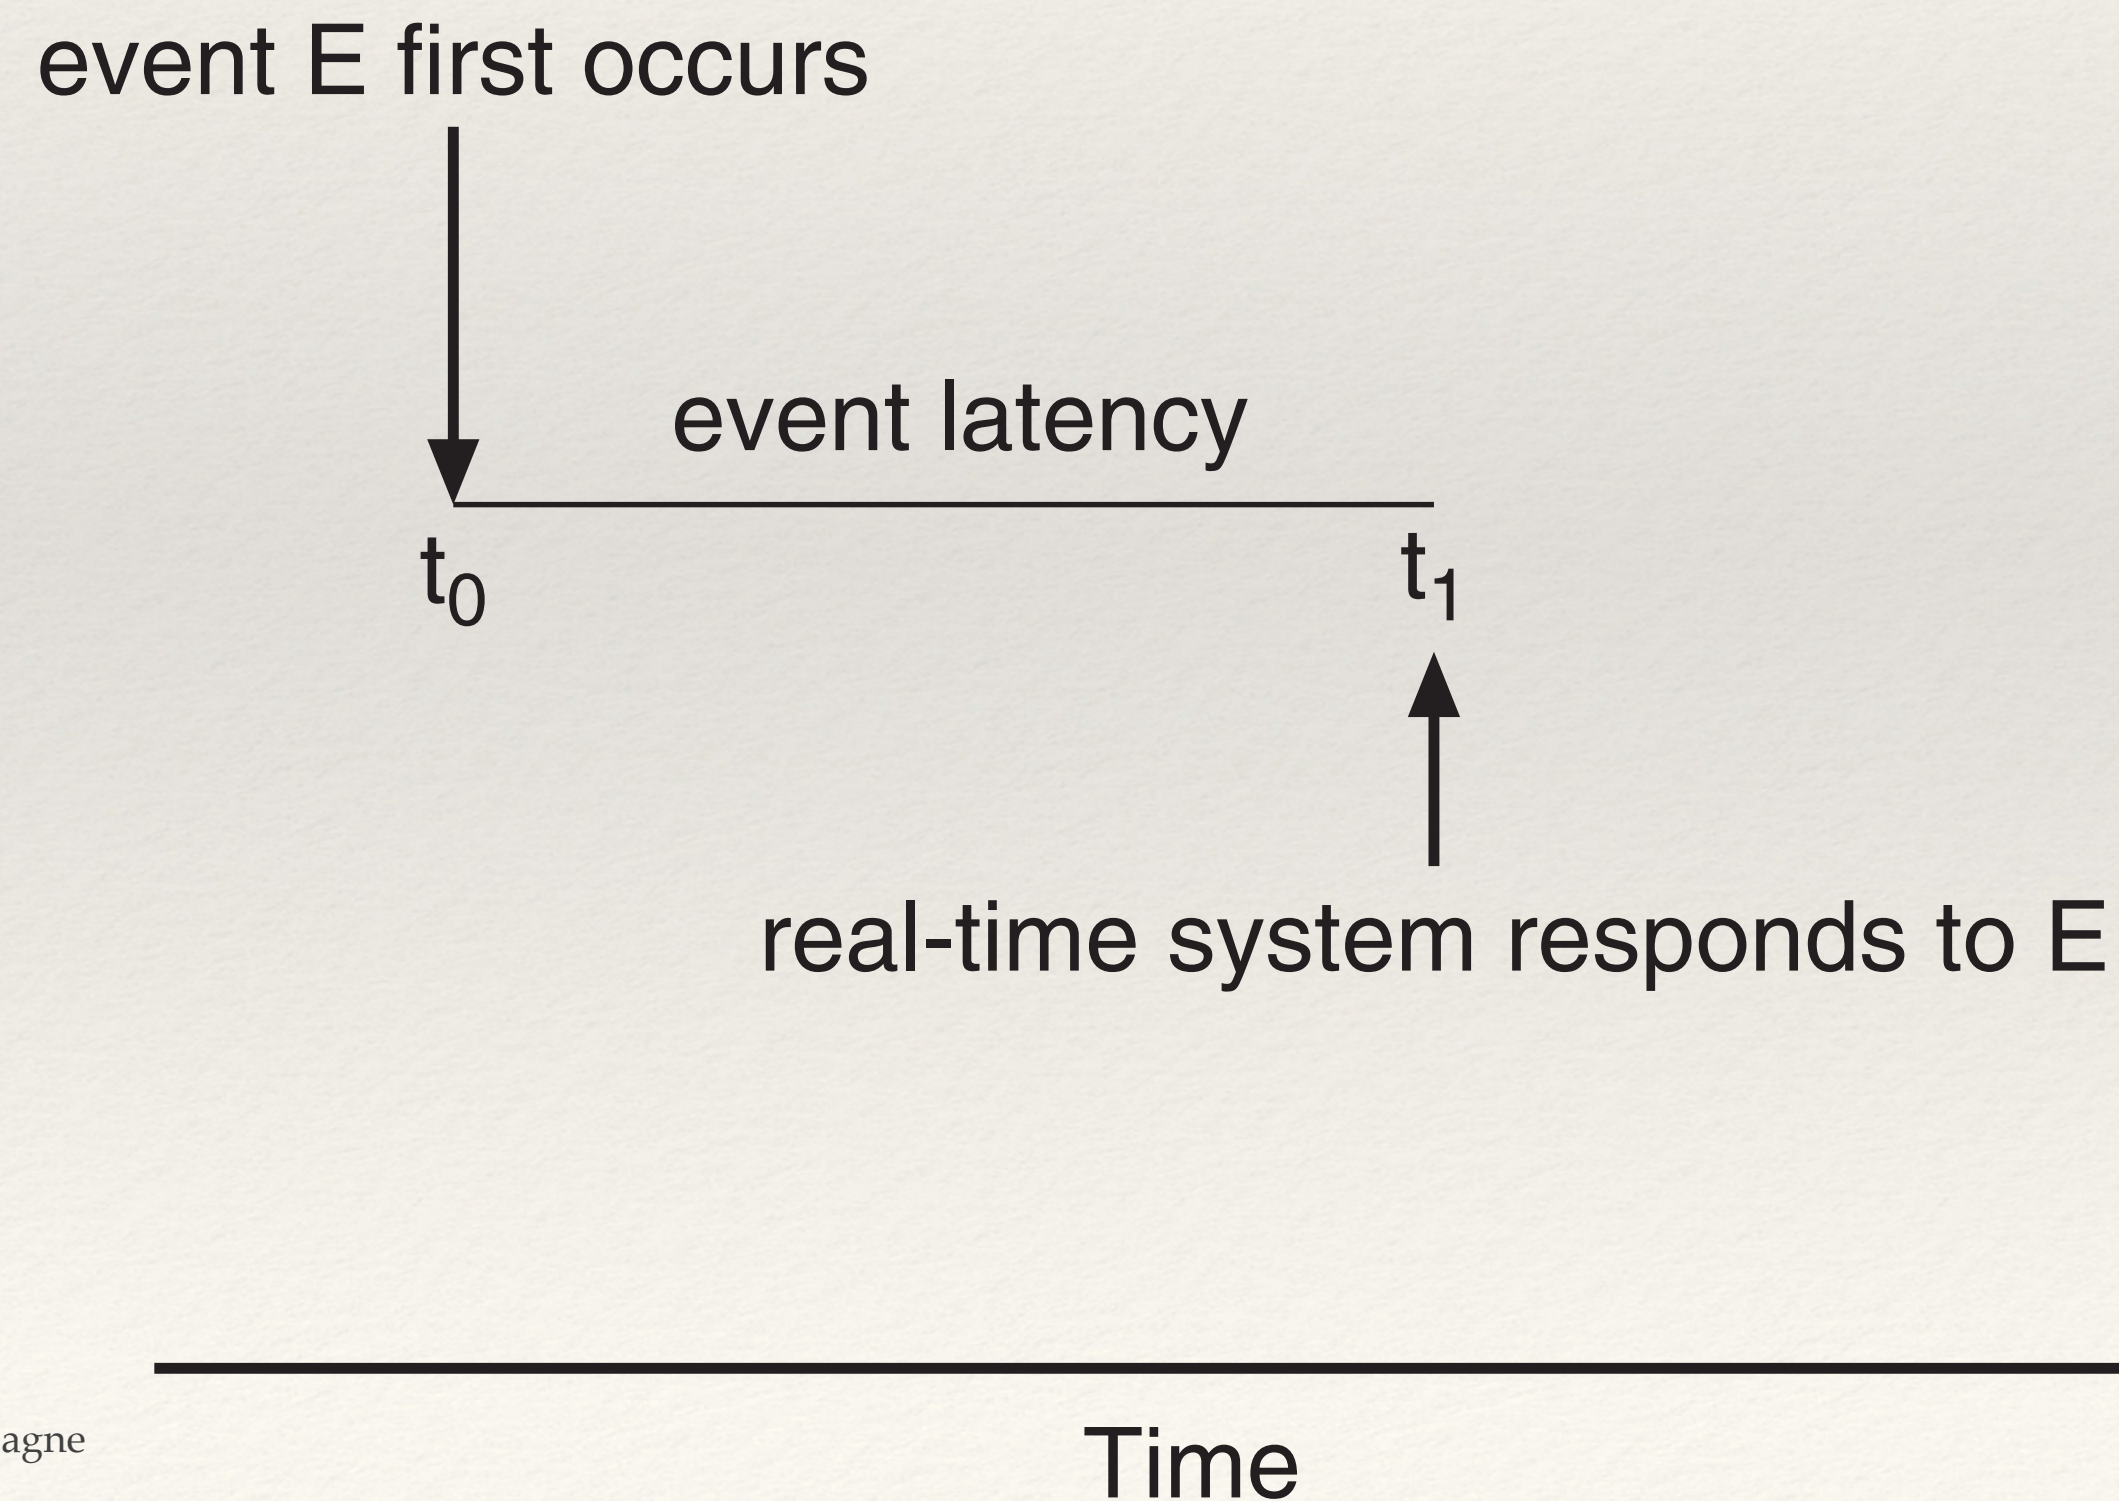


Real Time CPU Scheduling

- ❖ Soft real-time systems - no guarantee for critical real time process to be scheduled
- ❖ Hard real-time systems - stricter requirements, task must be serviced by it's deadline

Event Latency

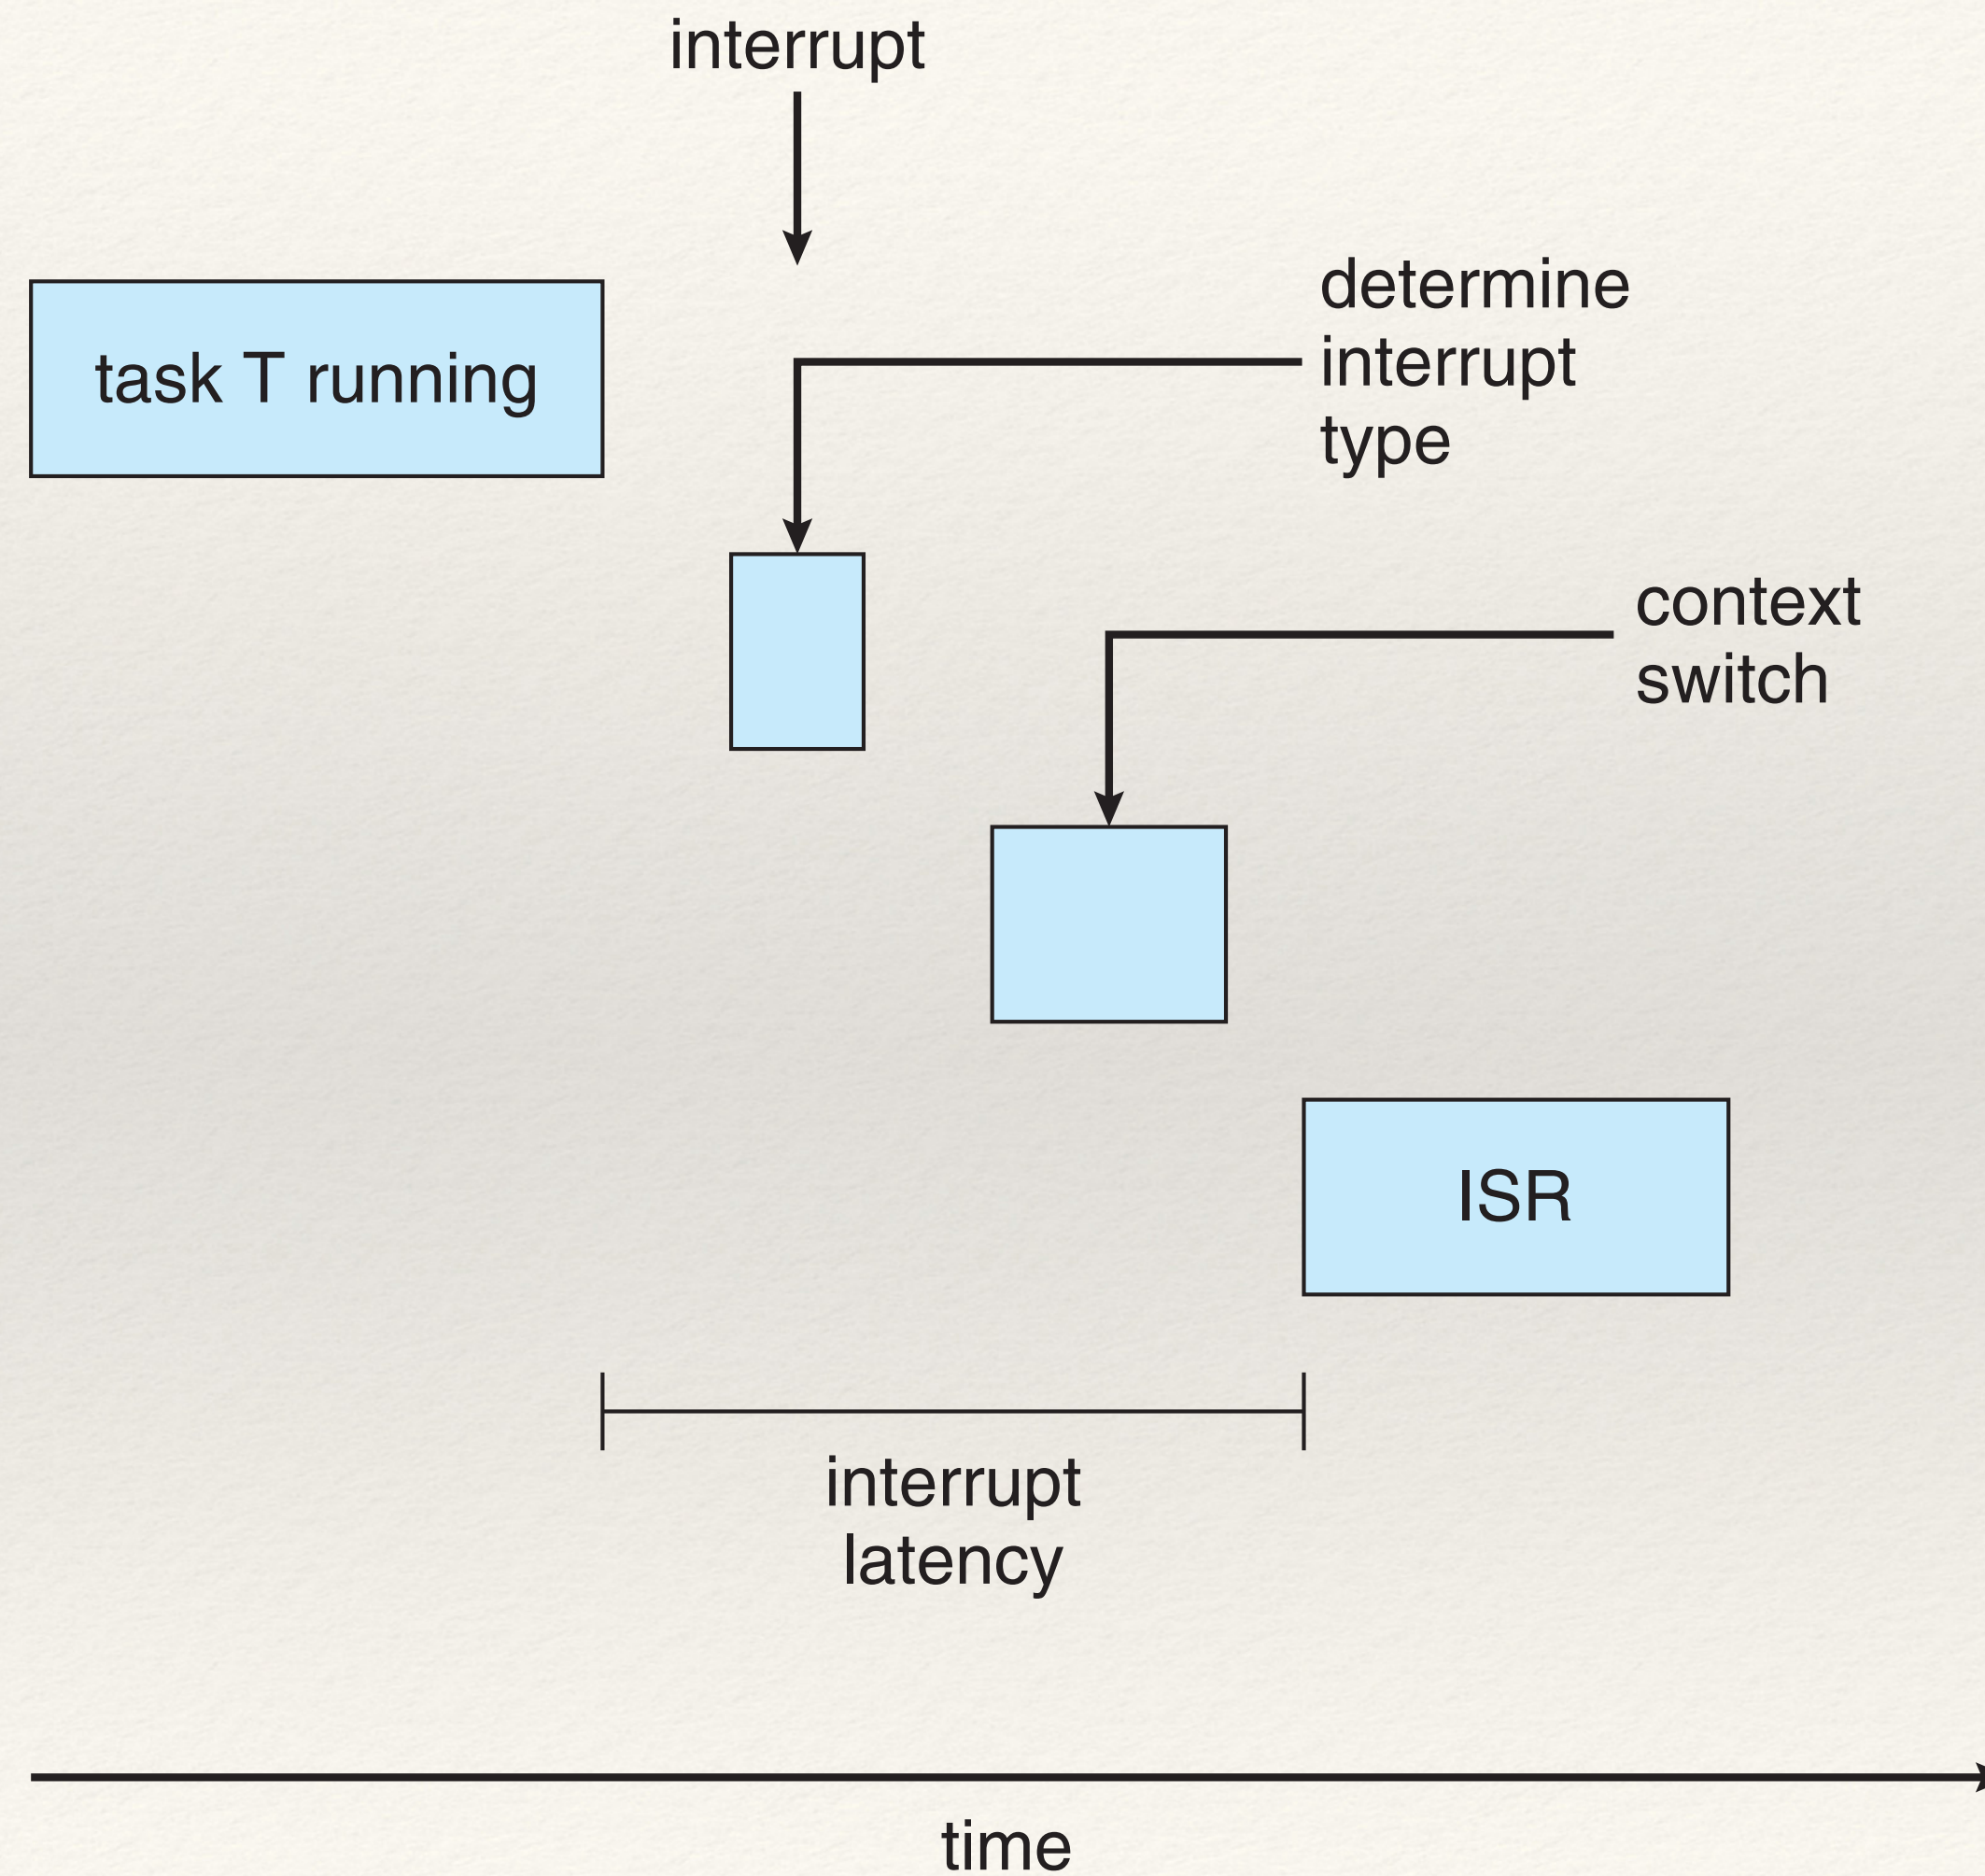
- ❖ Event Latency - amount of time that elapses from when an event occurs to when it's serviced



Types of Latency

- ❖ Interrupt Latency
- ❖ Dispatch Latency

Interrupt Latency

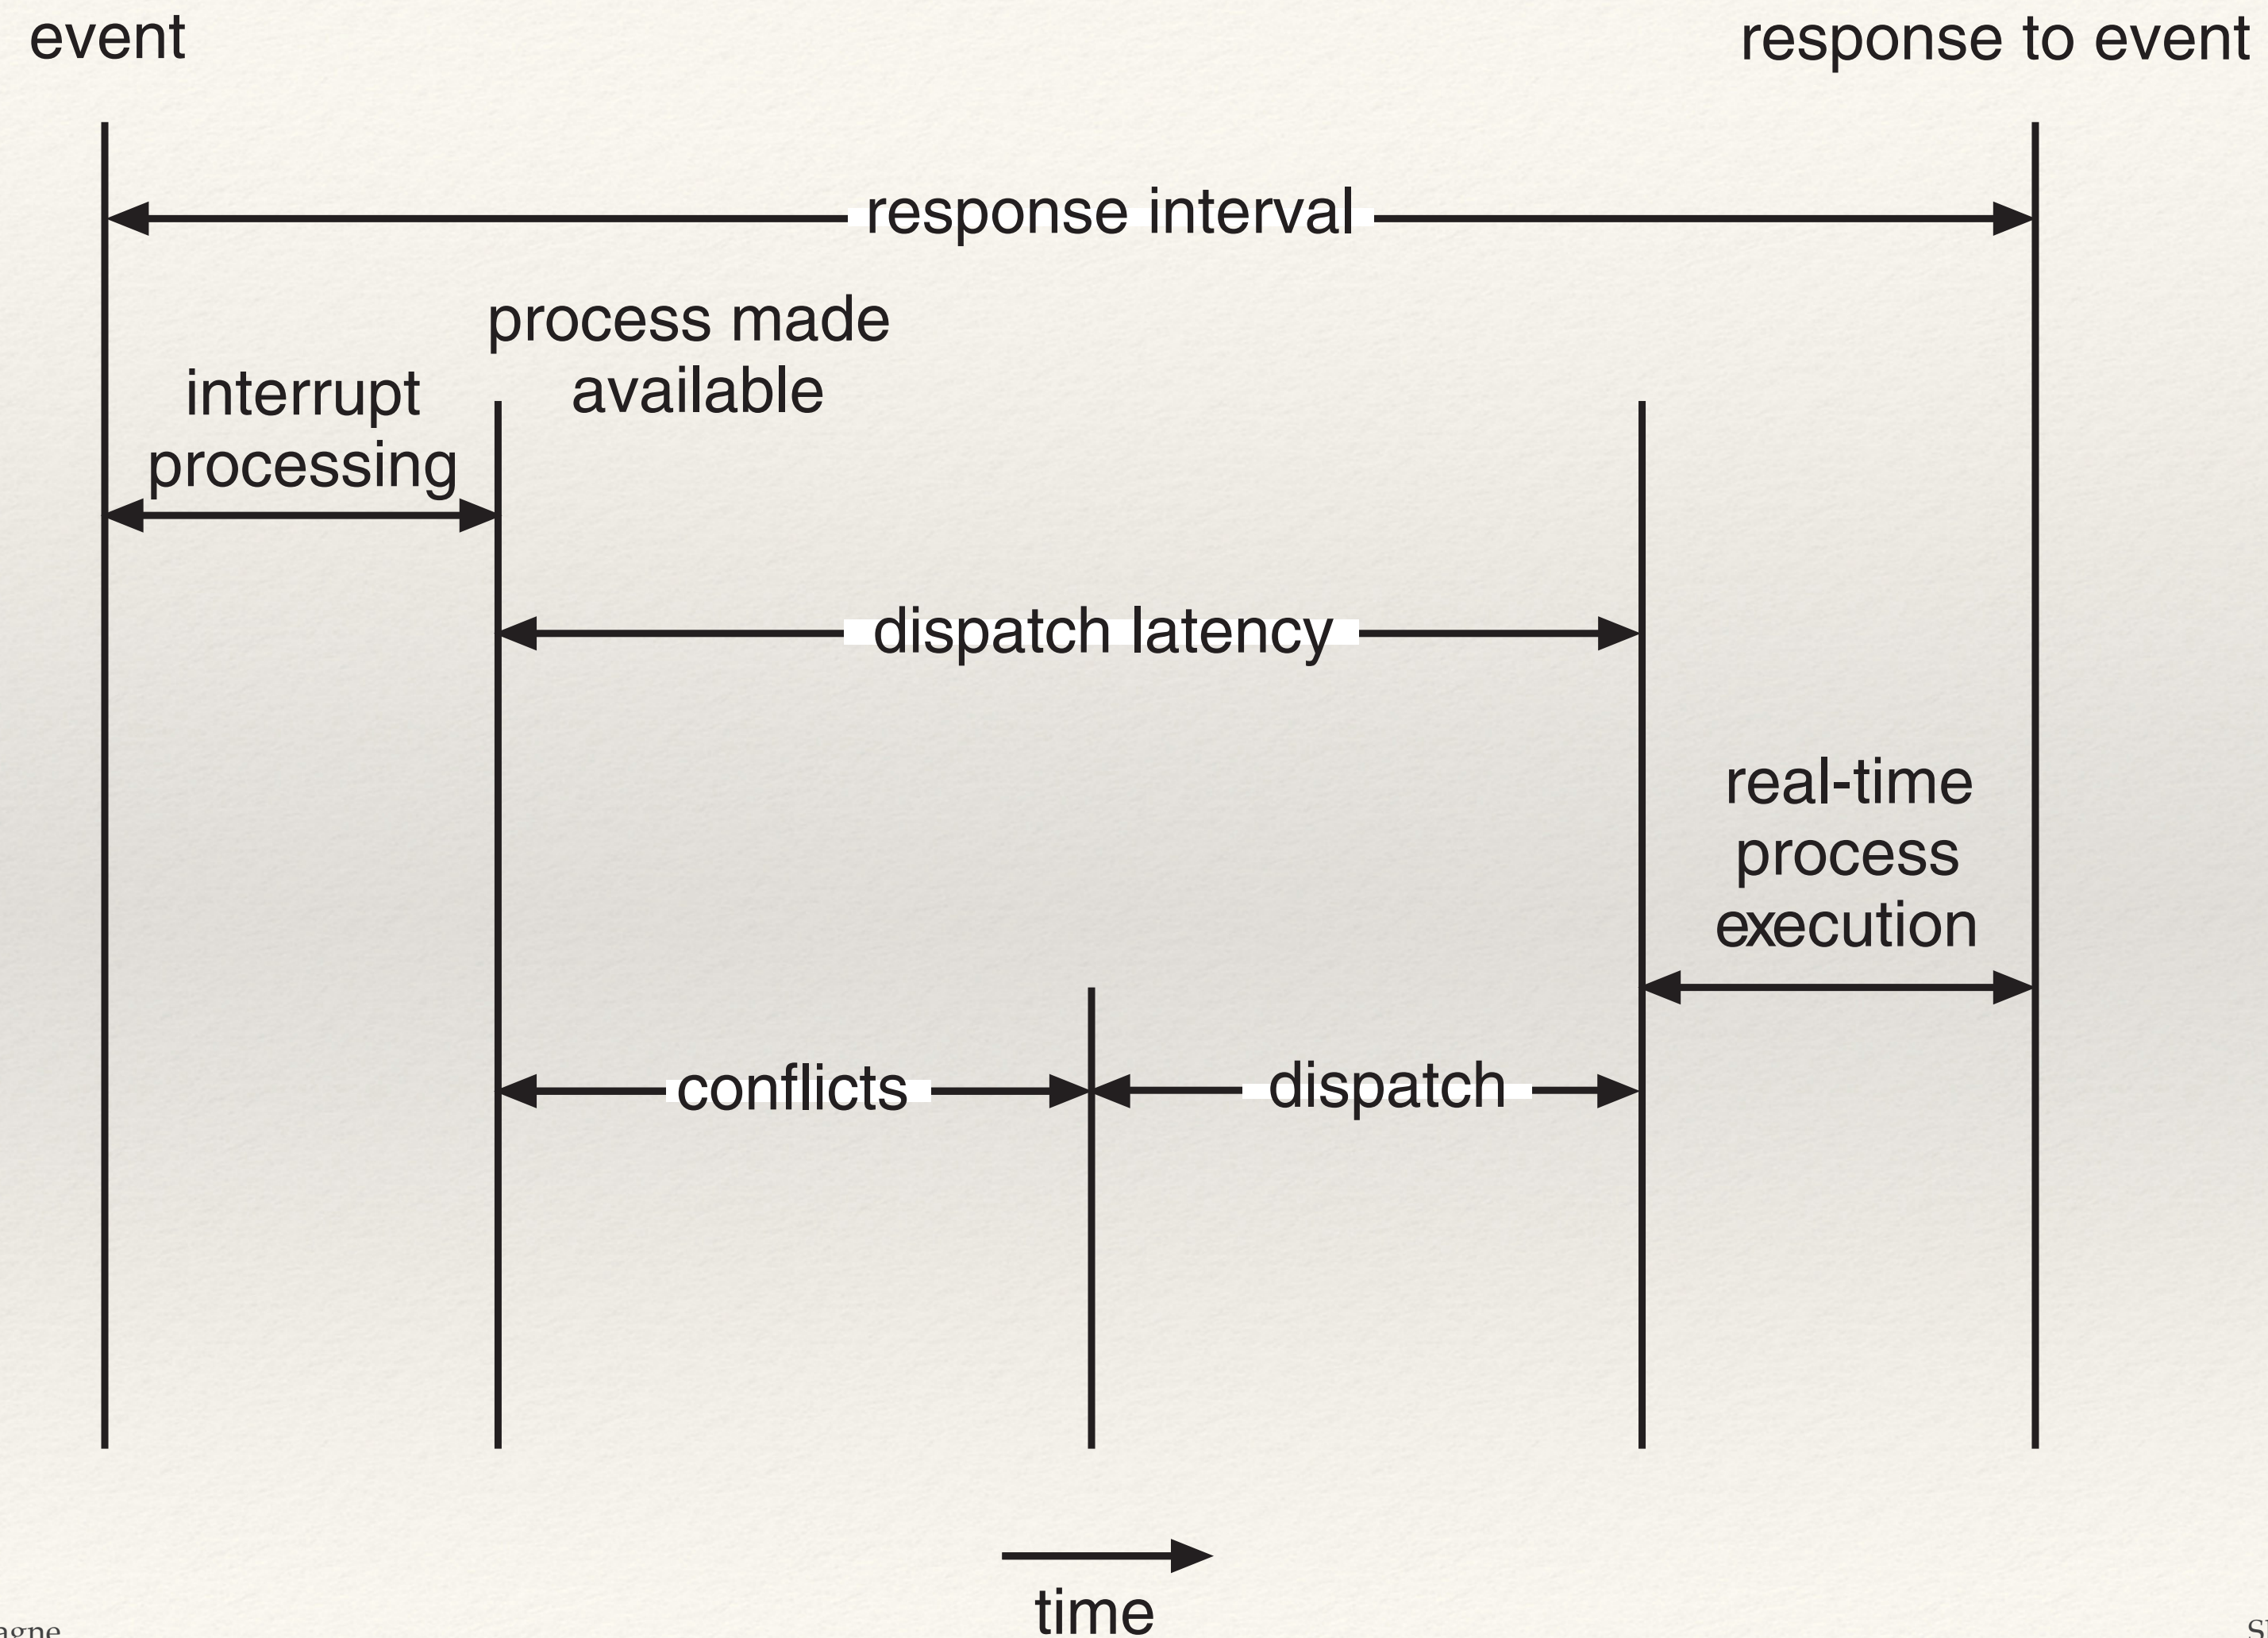


Interrupt Latency

- ❖ When an interrupt occurs
 - ❖ The OS must first complete the instruction it is executing and determine the type of interrupt that occurred.
 - ❖ Save the state of the current process
 - ❖ Service the interrupt using the Interrupt service routine (ISR)

Interrupt Latency is the time it takes to do the above tasks

Dispatch Latency



Dispatch Latency

- ❖ Scheduling dispatcher to stop one process and start another
- ❖ Conflict phase
 - ❖ Preemption of any process running in the kernel
 - ❖ Release of low priority processes of resources needed by a high-priority process

Priority based scheduling

Periodic processes require the CPU at certain intervals


The scheduler can guarantee the requirements or reject them (Admission-control algorithm)

p = period

d = deadline

t = fixed processing time

The rate of a period task is $1/p$



Rate-Monotonic Scheduling

- ❖ Used to schedule periodic tasks
- ❖ Assign a task a priority based on it's period
 - ❖ The shorter the period, the higher priority
 - ❖ The longer the period, the lower priority
- ❖ Assume each CPU burst is the same every period



Rate-Monotonic Example

- ❖ Process One: $p_1 = 50, t_1 = 20, d_1 = 50$
- ❖ Process Two: $p_2 = 100, t_2 = 35, d_2 = 100$

Rate-Monotonic Example

- ❖ Process One: $p_1 = 50, t_1 = 20, d_1 = 50$
 - ❖ Calculate CPU utilization $20/50 = 0.40$
- ❖ Process Two: $p_2 = 100, t_2 = 35, d_2 = 100$
 - ❖ Calculate CPU utilization $35/100 = .35$
- ❖ If the sum > 1 , not possible to schedule them on the same CPU.

Rate Monotonic Example



Earliest-Deadline-First Scheduling

- ❖ Schedule priorities dynamically
- ❖ Earlier the deadline, the higher the priority

Proportional Share Scheduling

- ❖ Scheduler allocates T shares among all applications
- ❖ Each application can receive N shares of time
 - ❖ This ensures that each application will have N/T of the total processor time
- ❖ Admission-control policy admits requests or denies request based on the number of shares that are available