

*CS 452 Operating Systems*

---

# Processes

Dr. Denton Bobeldyk



---

# History

---

- ❖ Early computers allowed only one program to execute at a time
  - ❖ Complete control of the system
  - ❖ Access to all the systems resources (memory, disk, etc)
- ❖ Modern computers allow for more than one program (process) to operate at a time




---

# Process

---

- ❖ Definition: A program in execution
  - ❖ A program itself is a passive entity
  - ❖ A process is an active entity





**Figure 3.1** Process in memory.

# Process

**Stack:** Contains temporary data such as function parameters, return addresses, local variables

**Heap:** memory dynamically allocated during process run time

**Data Section:** Global variables

**Text section:** Program code



---

# Multiple Processes

---

- ❖ A single program could run as multiple processes
  - ❖ For example, running two copies of Word or a command line app
  - ❖ These processes aren't associated with each other and are separate copies



---

# Process States

---

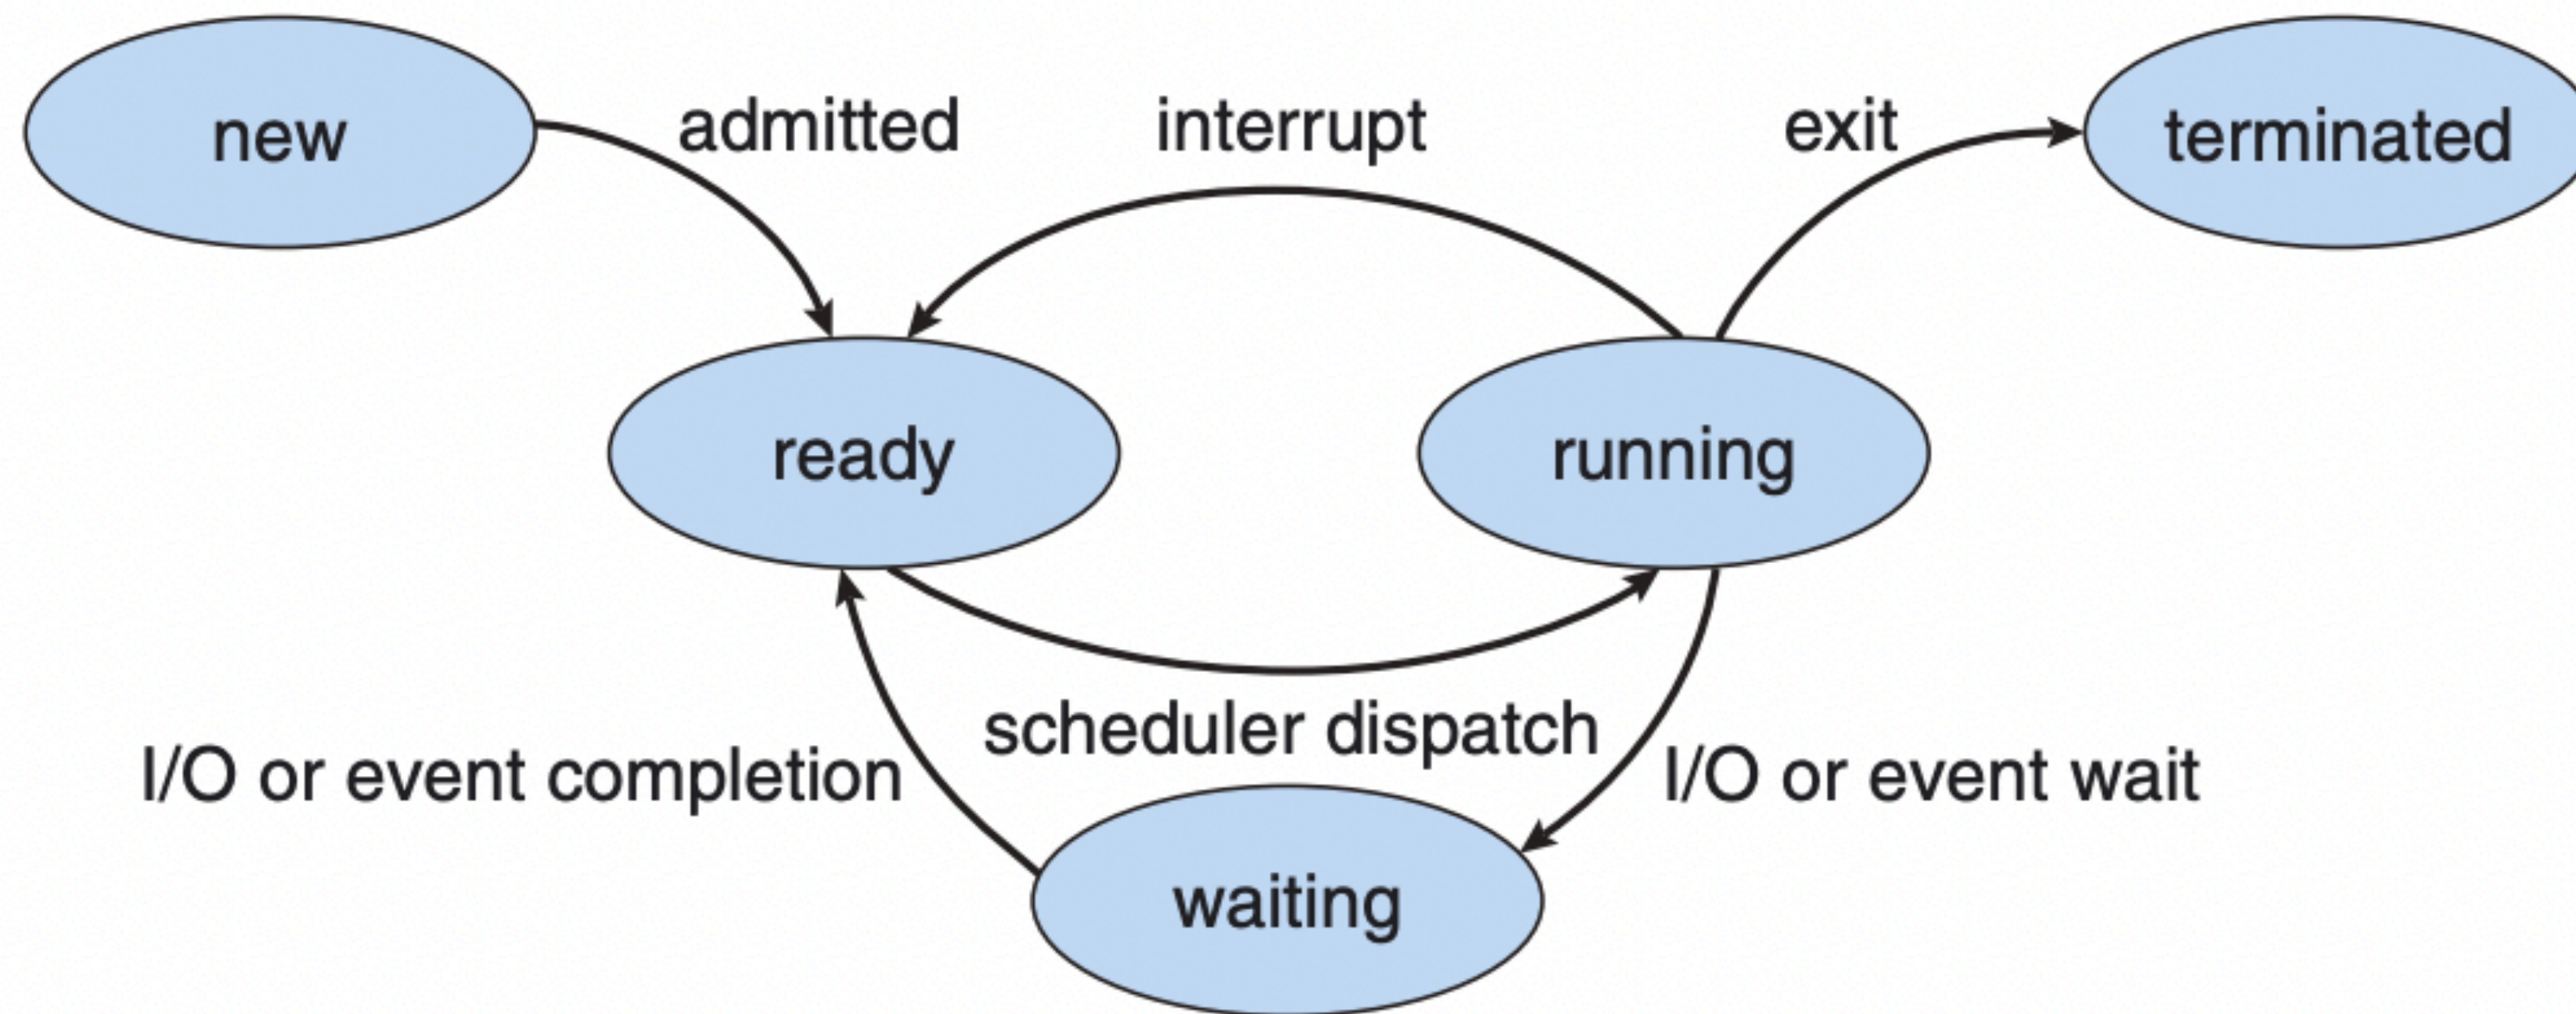
- ❖ New - Process is being created
- ❖ Running - Instructions are executed
- ❖ Waiting - The process is waiting for some event to occur
- ❖ Ready - The process is waiting to be assigned to a processor
- ❖ Terminated - The process has finished execution



Image from: [https://en.wikipedia.org/wiki/Terminator\\_\(character\)](https://en.wikipedia.org/wiki/Terminator_(character))



# Process States



**Figure 3.2** Diagram of process state.



---

# Process Control Block

---

- ❖ Process Control Block - each process represented in the operating system by a process control block



---

# Process Control Block

---

- ❖ Process State - {New, Ready, Running, Waiting, Terminated}
- ❖ Program Counter - indicates the address of the next instruction to be executed
- ❖ CPU registers - saves state of cpu registers for when an interrupt occurs. May include accumulators, index registers, stack pointers, ...



---

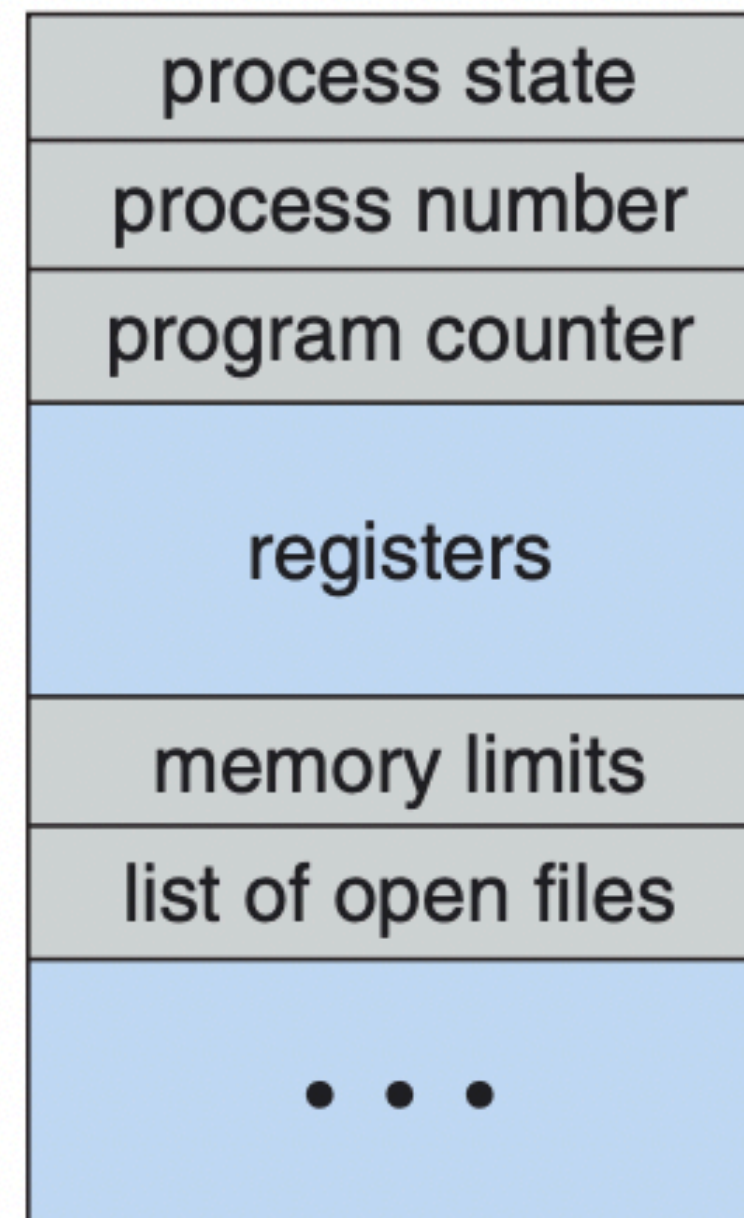
# Process Control Block

---

- ❖ CPU-scheduling information - process priority, pointers to scheduling queues, ...
- ❖ Memory-management information - base and limit registers, page tables, segment tables (covered in chapter 7)
- ❖ Accounting Information - CPU time used, time limits, account numbers, process numbers
- ❖ I/O status information - list of I/O devices allocated, list of open files, ...

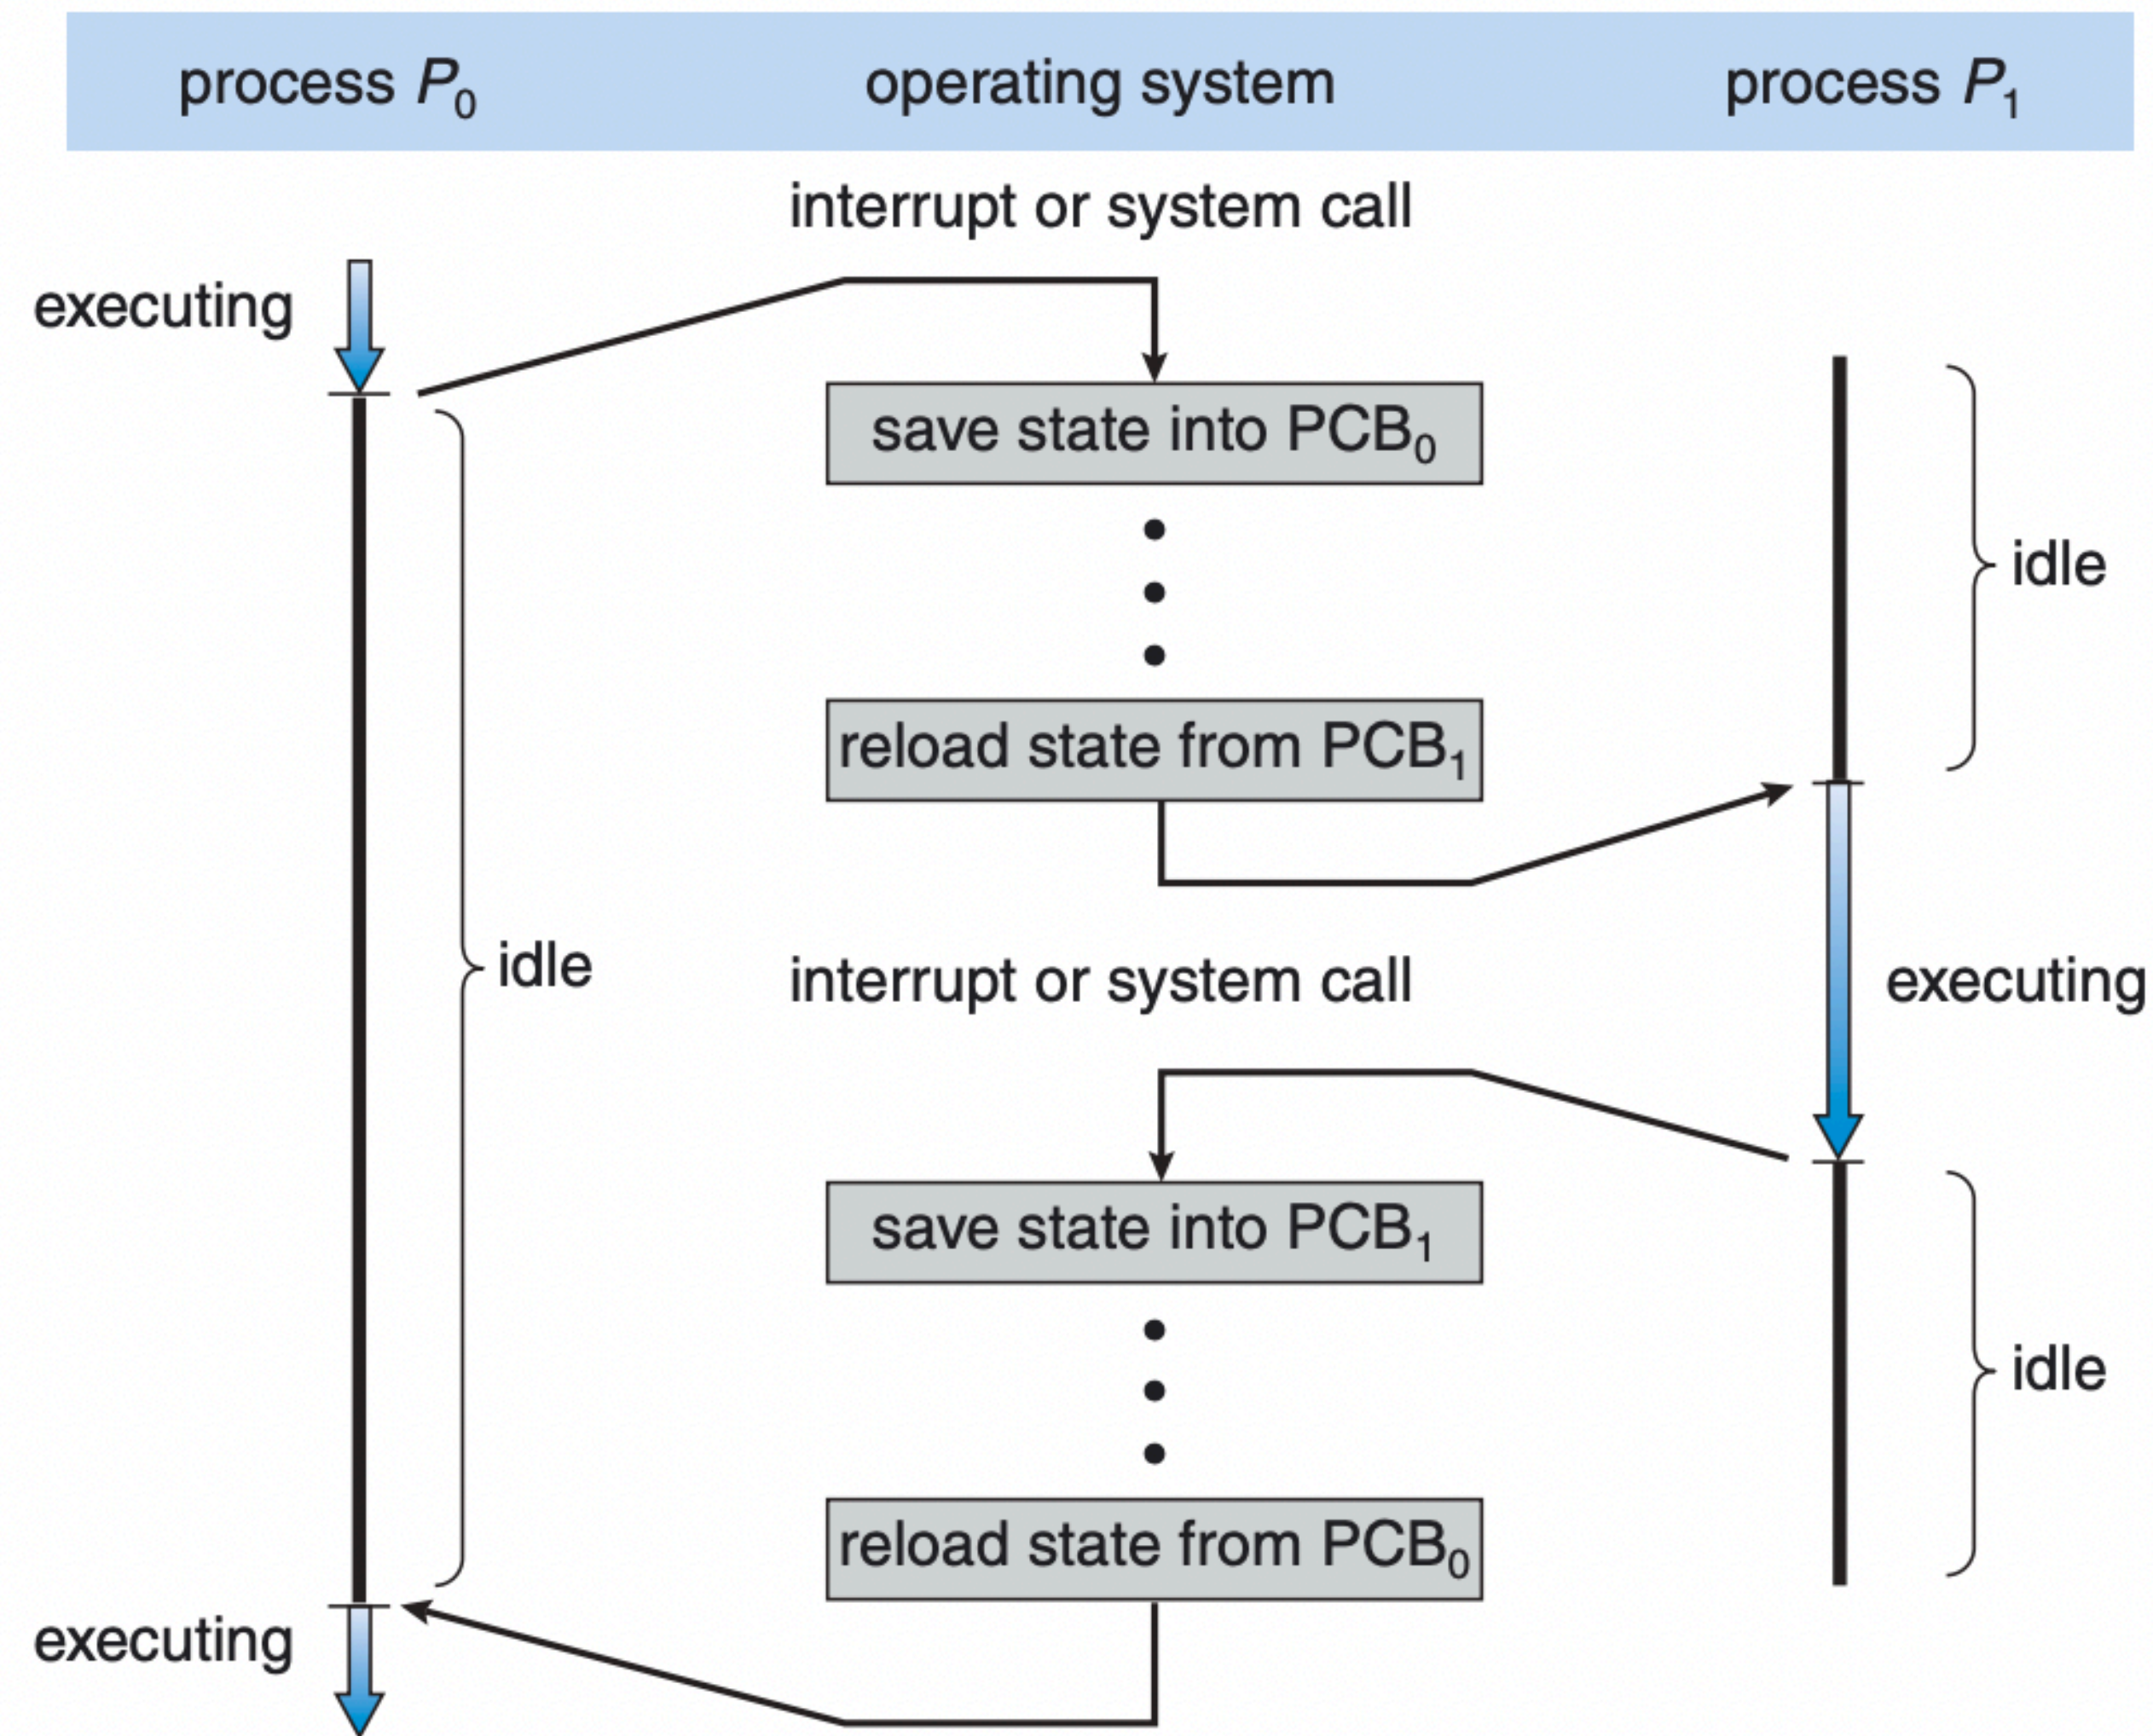


# Process Control Block



**Figure 3.3** Process control block (PCB).





**Figure 3.4** Diagram showing CPU switch from process to process.



---

# Threads

---

- ❖ A single thread allows a process to perform only one task at a time
- ❖ A multithreaded process allows a GUI to be responsive and to perform a task at the same time.
- ❖ Threads will be covered in Chapter 4



---

# Process Scheduling

---

- ❖ Multiprogramming - goal is to have a process running at all times
- ❖ Time Sharing - switch processes so frequently that users are able to interact with each



---

# Process Scheduling

---

- ❖ Multiprogramming - goal is to have a process running at all times
- ❖ Time Sharing - switch processes so frequently that users are able to interact with each
- ❖ Process scheduler helps us accomplish this!



---

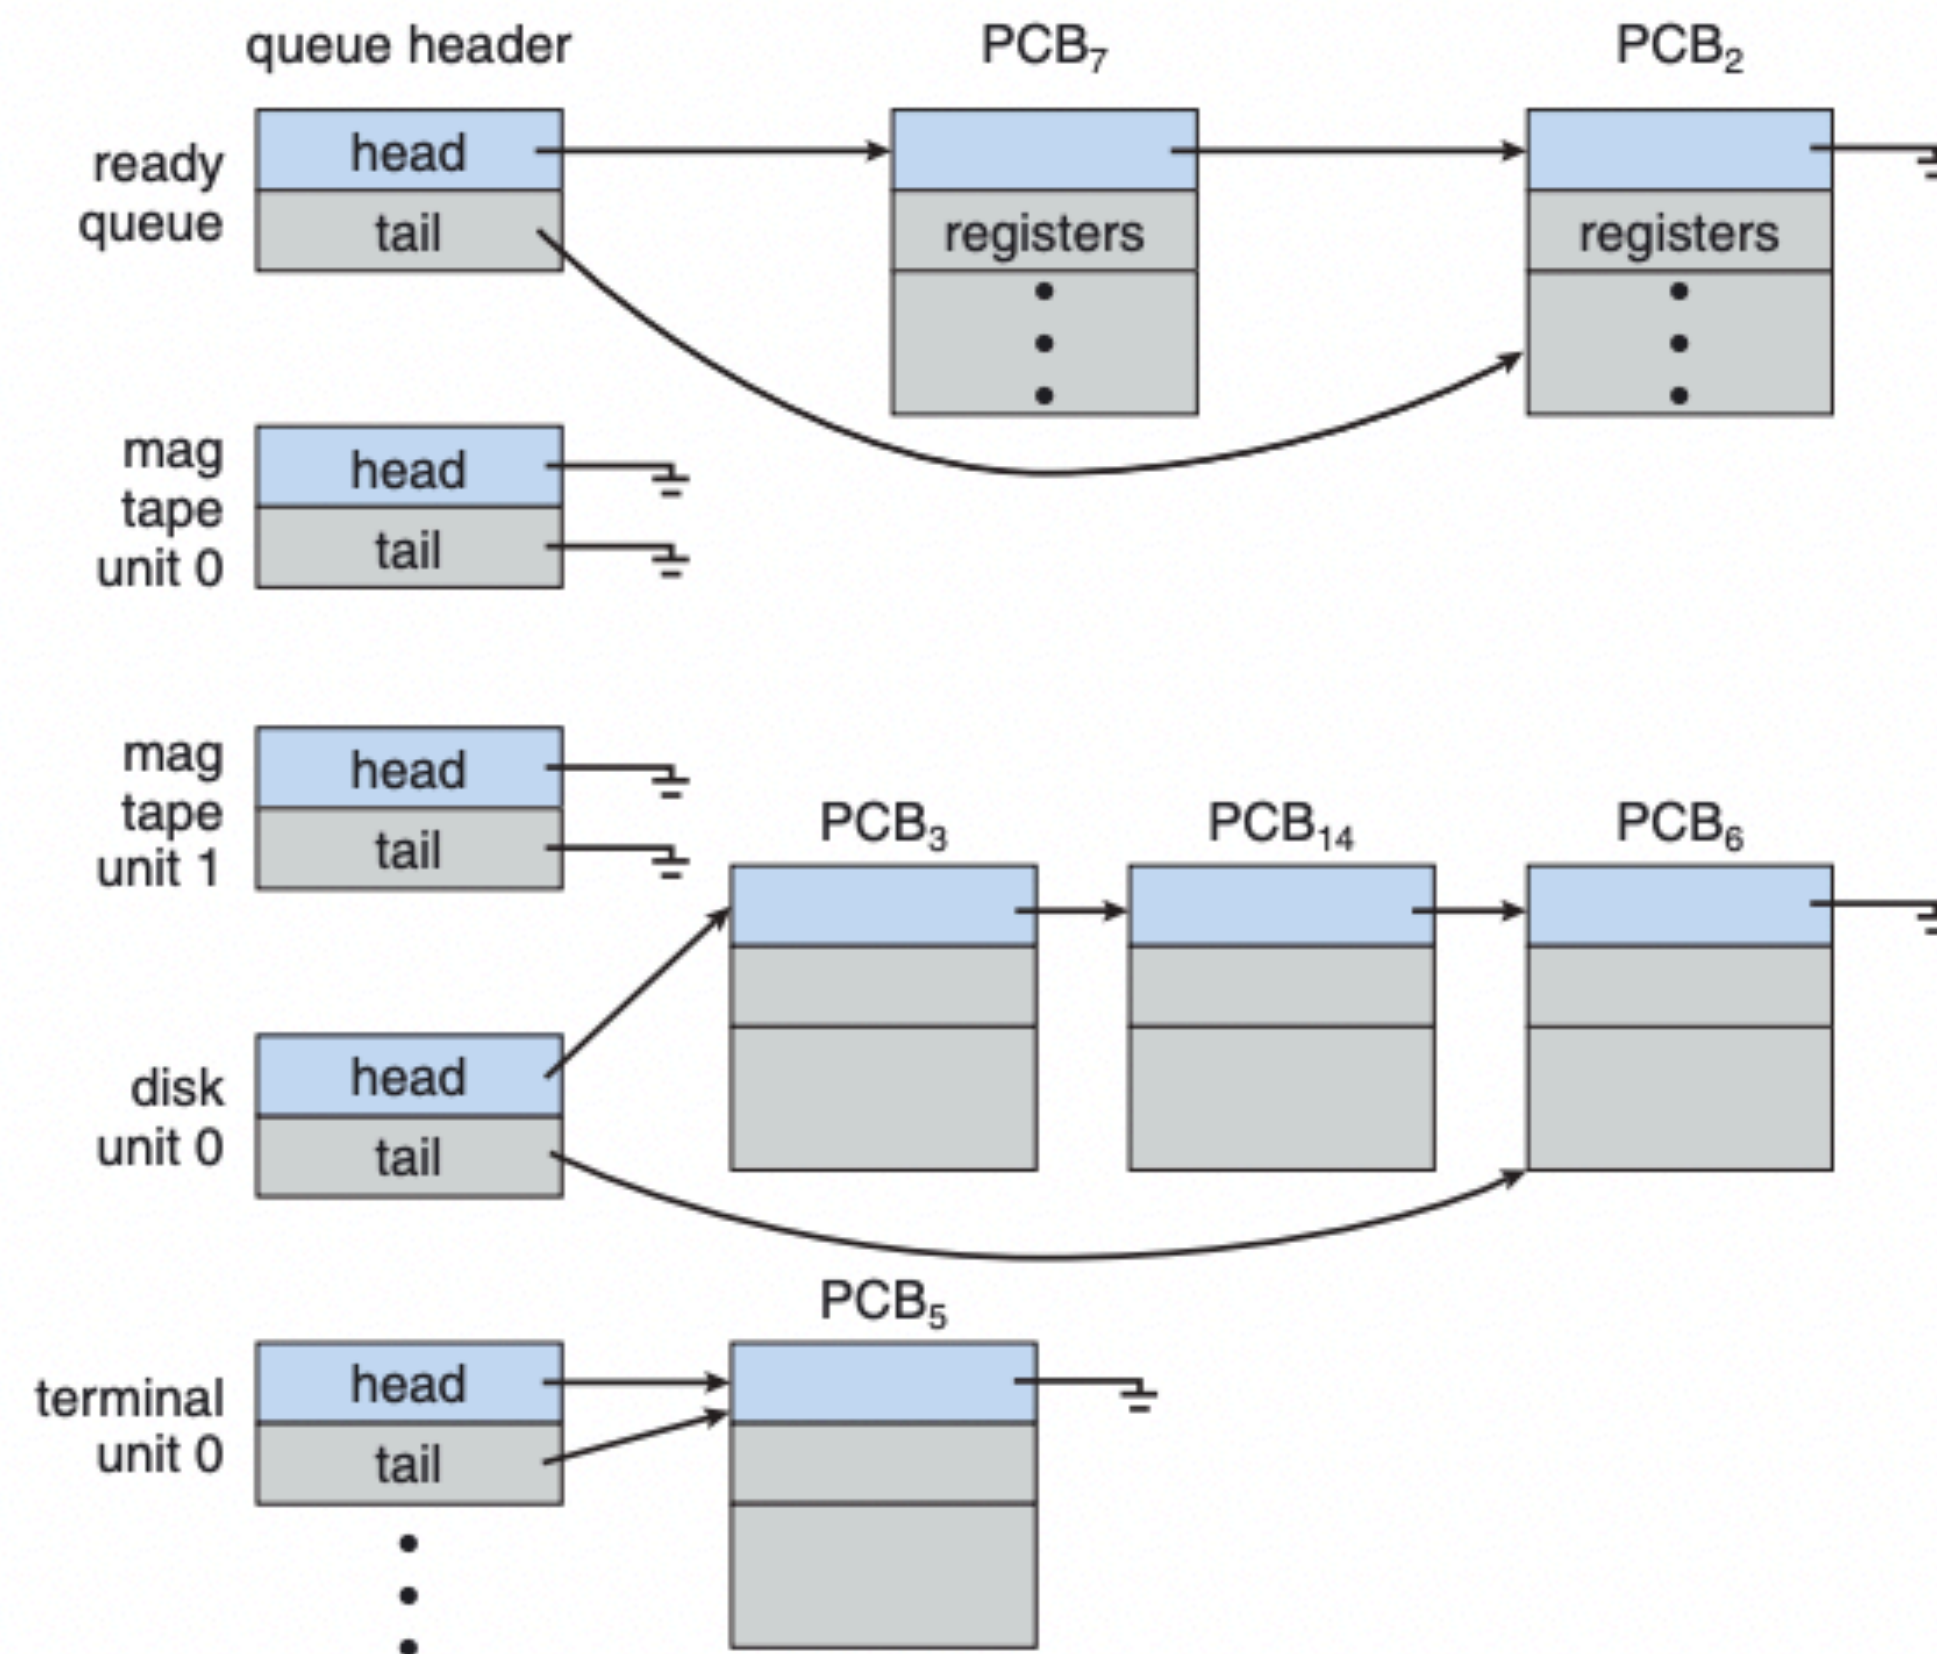
# Scheduling Queues

---

- ❖ Job Queue - all processes in a system
- ❖ Ready Queue - processes residing in main memory, ready and waiting to execute
- ❖ Device Queue - processes waiting for an I/O device. Each device has it's own queue

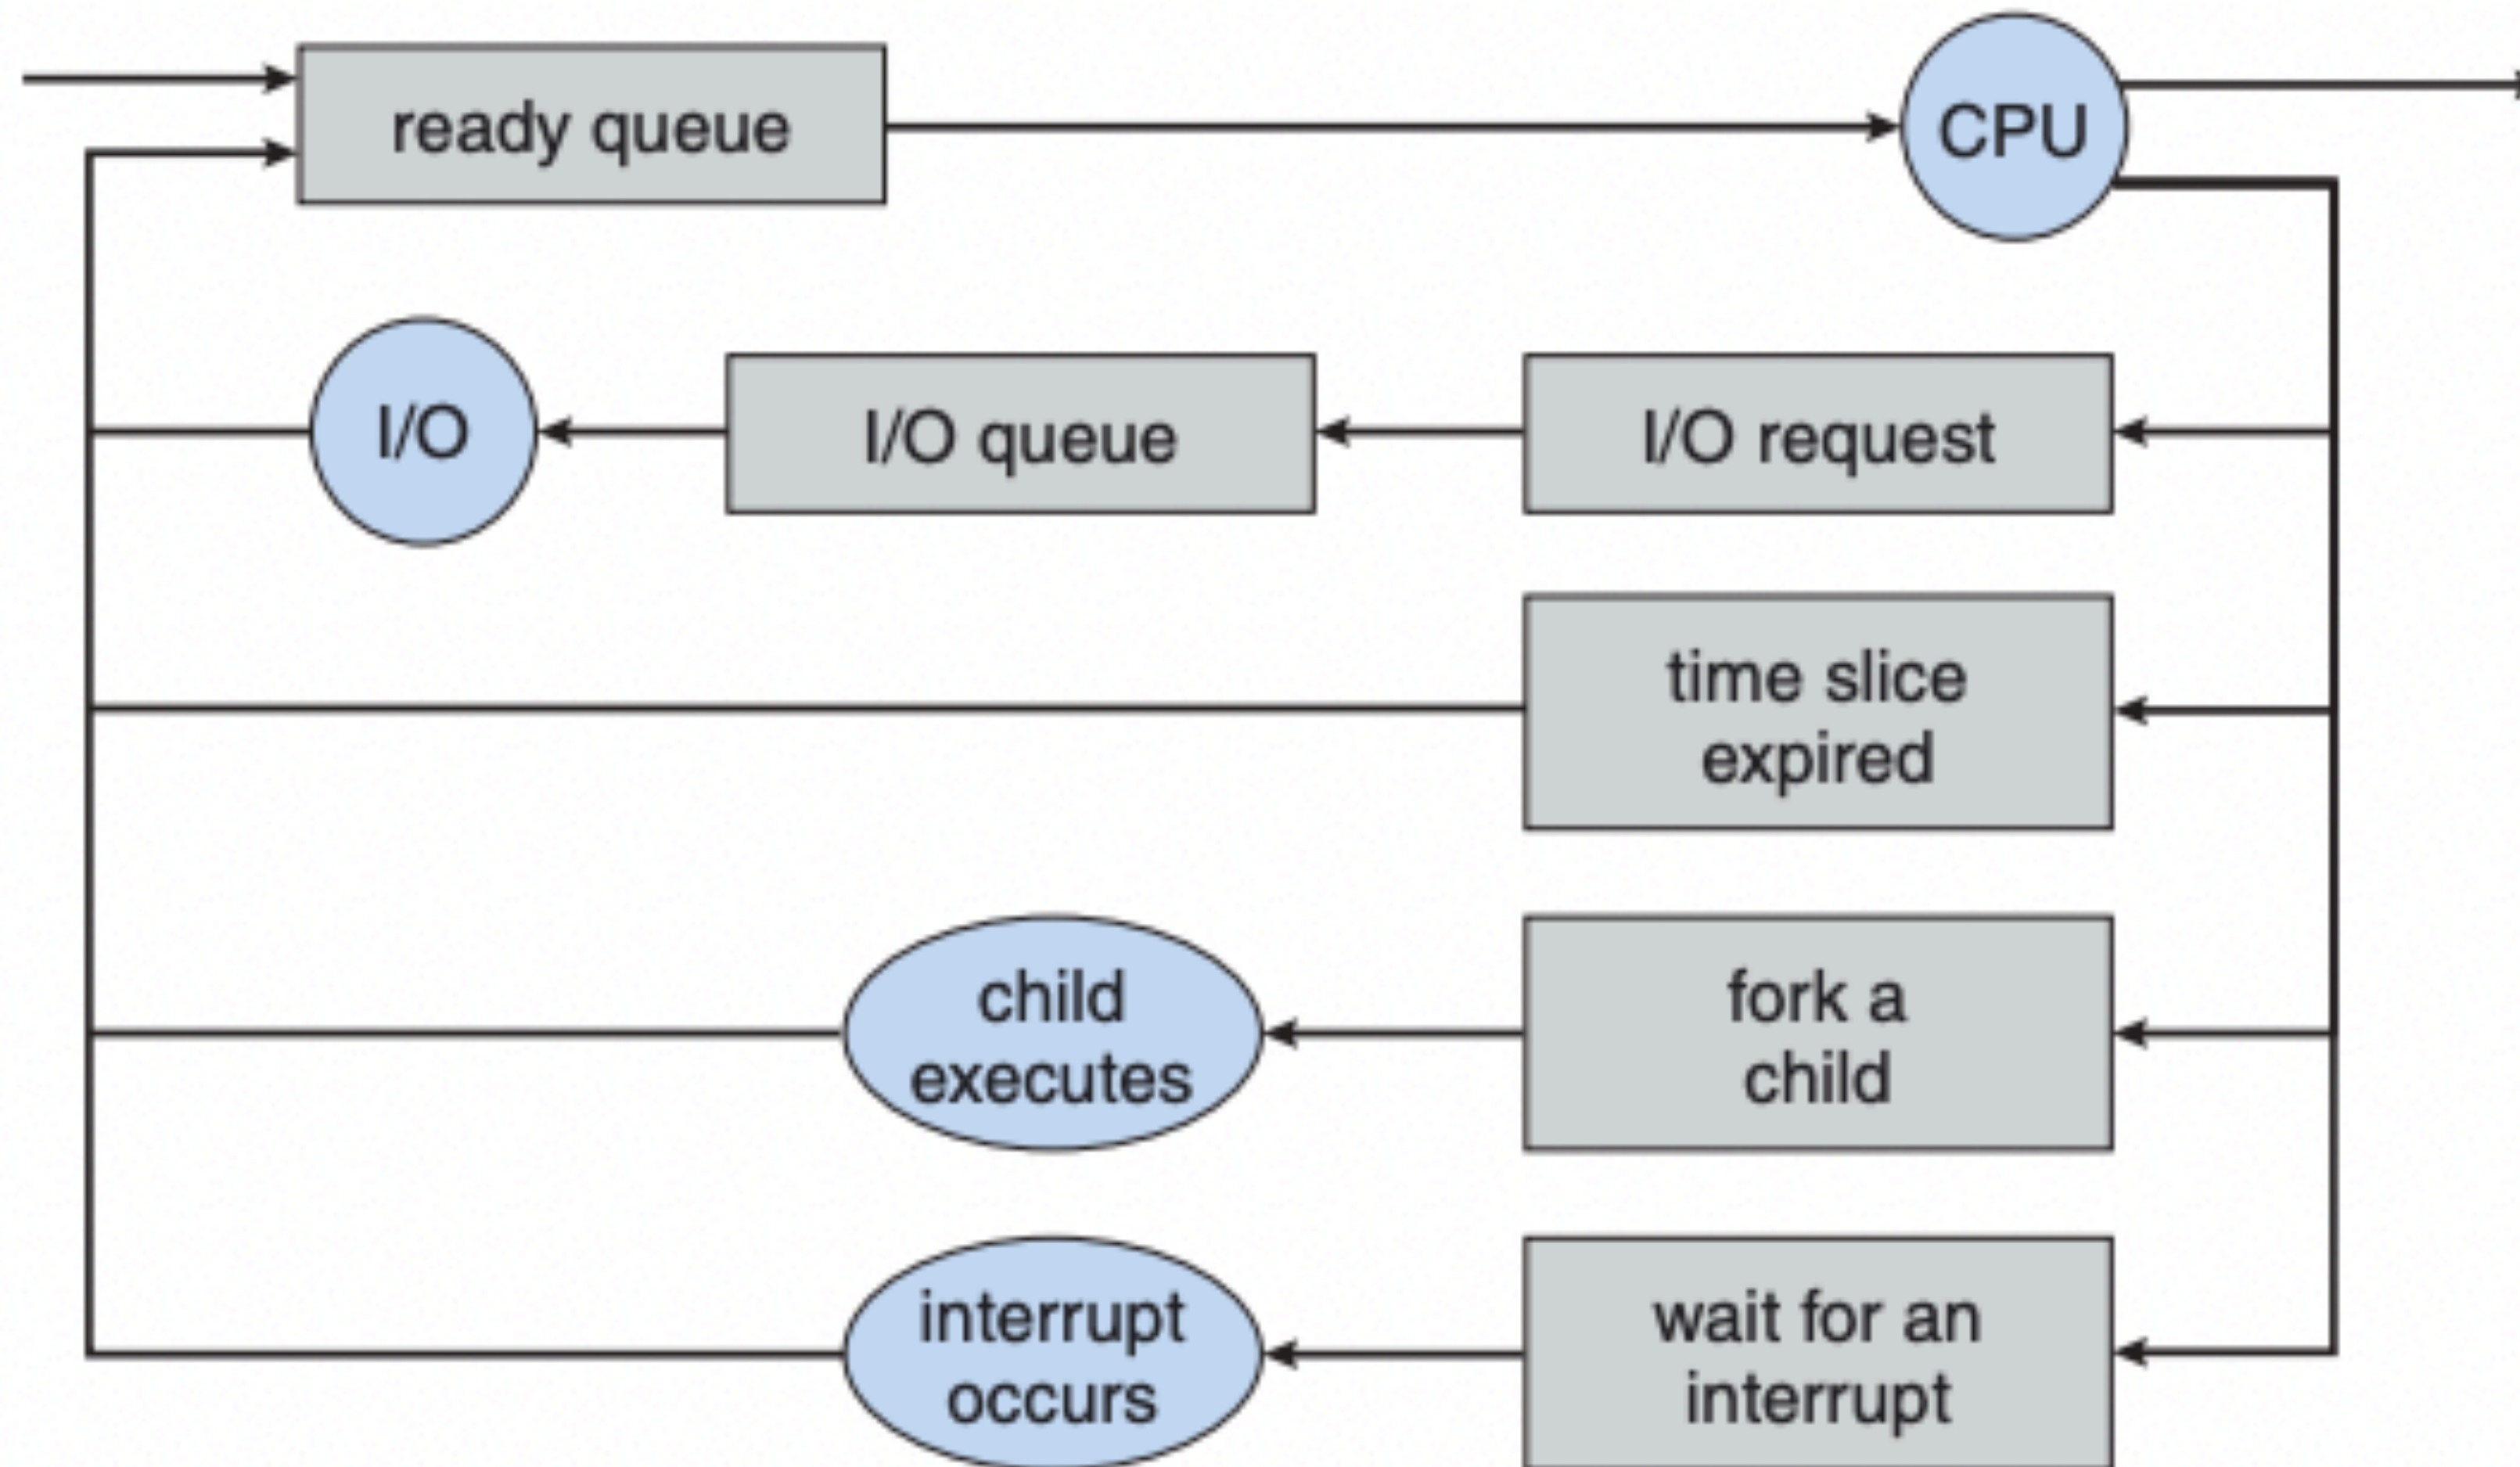


# Scheduling Queues



**Figure 3.5** The ready queue and various I/O device queues.





**Figure 3.6** Queueing-diagram representation of process scheduling.



---

# Schedulers

---

- ❖ Long-term scheduler or job scheduler - selects processes from a pool to be executed. The processes selected are loaded into memory
- ❖ Short-term scheduler - selects from processes already in memory



---

# Schedulers

---

- ❖ Long-term scheduler or job scheduler - runs fairly infrequently, can take longer times to make decisions, not available on all computers
- ❖ Short-term scheduler - runs frequently, needs to make decisions faster



---

# Processes

---

- ❖ I/O bound - spends more time doing I/O than doing computations
- ❖ CPU bound - spends more time doing computations than waiting for I/O



---

# Context Switch

---

- ❖ Interrupt occurs, need to save context of current process.
- ❖ Context is saved as a Process Control Block (PCB)
- ❖ Switching the CPU to another process, and saving the state of the current process is known as a **context switch**.



---

# Context Switch

---

- ❖ What type of hardware might help speed up a context switch?



---

# Context Switch

---

- ❖ What type of hardware might help speed up a context switch?
  - ❖ Hardware support
    - ❖ Multiple sets of registers and the context switch is simply changing a pointer to the running process.



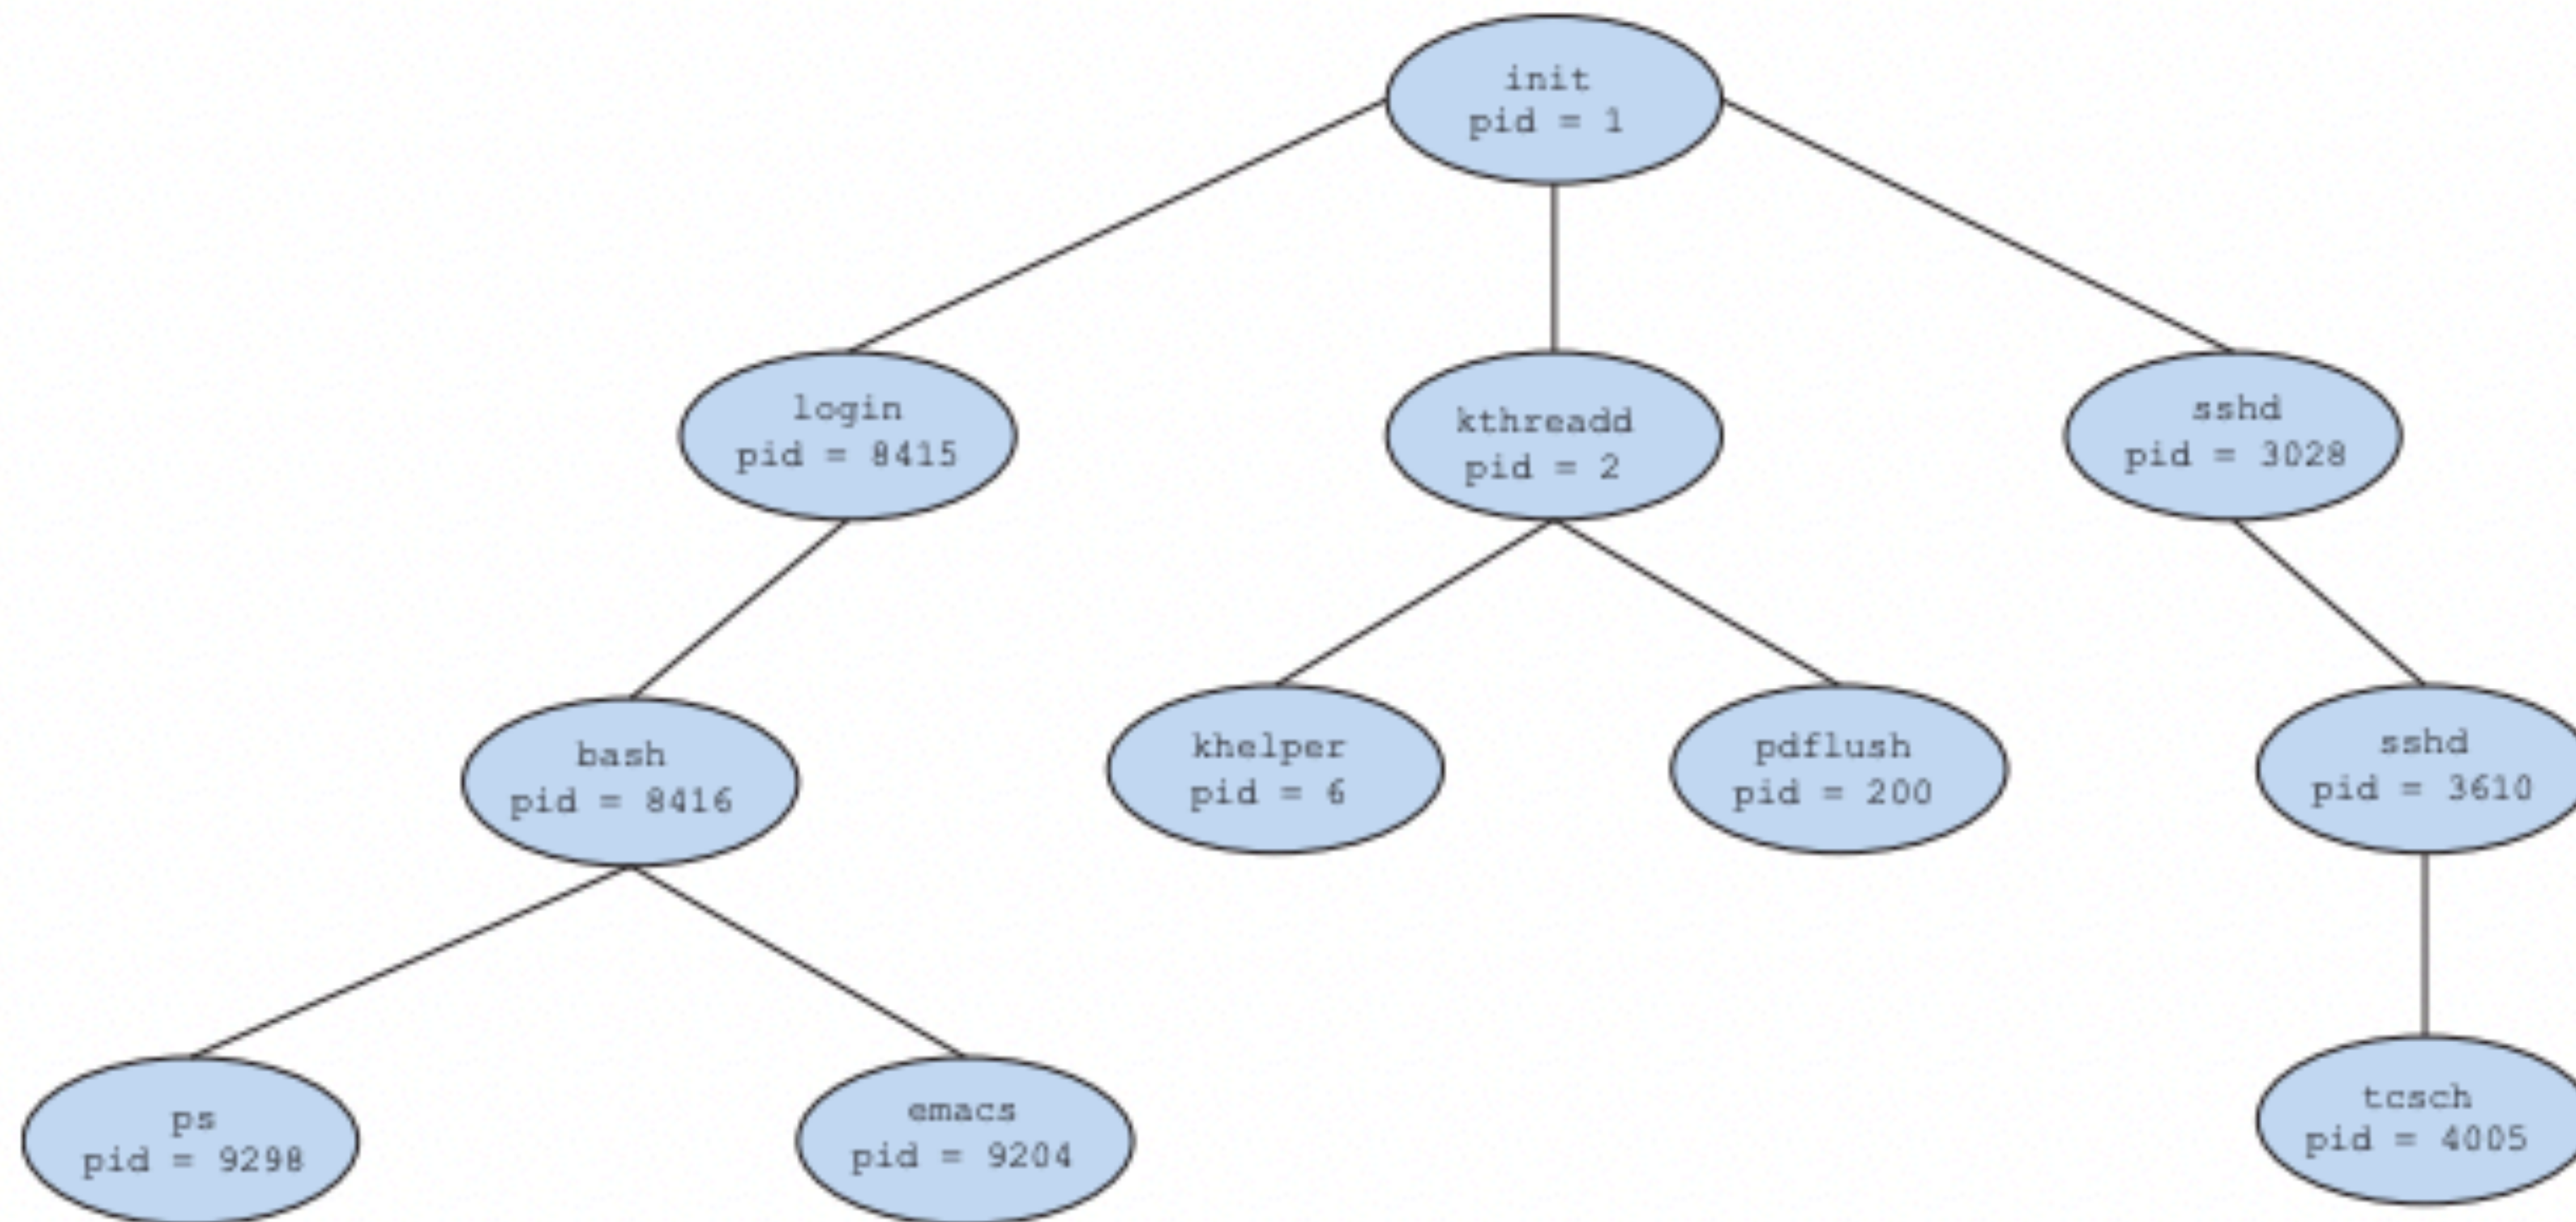
---

# Process Creation

---

- ❖ A process (sometimes called a task) can create another process.
- ❖ Parent process - the creating process
- ❖ Child process - the process created
- ❖ Tree - parent and its children





**Figure 3.8** A tree of processes on a typical Linux system.



---

# Process Creation

---

- ❖ Child process
  - ❖ Access resources (memory, files, I/O devices, CPU time) directly from the Operating System
  - ❖ Access resources partitioned by the parent



---

# Process Creation

---

- ❖ What happens to the parent when it spawns a child?
  - ❖ Wait for child to execute and terminate?
  - ❖ Continue to execute concurrently with its children



---

# Process Creation

---

- ❖ Child process can be:
  - ❖ Duplicate of the parent process (same program / data)
  - ❖ New program loaded into it (using exec)
    - ❖ Replaces process memory space with the new program



# Process Creation In-Class Assignment

- ❖ Create a new process using the fork method
- ❖ To complete after the lecture (time permitting)

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>
#include <sys/wait.h> /*textbook doesn't have this, wait doesn't work without it*/

int main(void)
{

    pid_t pid;
    pid_t childProcessID;
    pid_t parentProcessID;

    parentProcessID = getpid();

    printf("\n[%d]This gets run once\n", getpid());

    /* fork a child process */
    childProcessID = fork();

    if(childProcessID < 0){ /* error occurred */
        printf("Fork Failed\n");
        return 1;
    }
}
```



---

# Process Creation

---

- ❖ How many children can you fork from a process?



---

# Process Creation

---

- ❖ How many children can you fork from a process?
  - ❖ Limited based on the amount of memory, maximum value for pid.
  - ❖ On a 64 bit system could be ~4 million, but typically related to resources



---

# Process Termination

---

- ❖ A process terminates when it finishes executing its final statement and asks the OS to delete it by using the `exit()` system call.
- ❖ A return value can be sent to its parent via the `wait()` system call
- ❖ All system resources are deallocated.



---

# Process Termination

---

- ❖ Parent typically is the only one that can terminate a child process
- ❖ Prevent users from terminating other users' processes



---

# Process Termination

---

- ❖ Possible reasons for child termination:
  - ❖ Child has exceeded usage of resources allocated
  - ❖ Task assigned to the child is no longer required
  - ❖ The parent is exiting and the OS does not allow the child to continue if it's parent is terminated
    - ❖ Cascading termination



---

# Process Termination

---

- ❖ Parent process may wait for a child process using the wait() system call

```
pid_t pid;
```

```
int status;
```

```
pid = wait(&status)
```



---

# Process Termination

---

- ❖ A process who has terminated, but the parent has not yet called wait() is called a zombie process
- ❖ Once the parent calls wait() the process identifier and **zombie** process and it's entry in the process table is released
- ❖ If the parent did not invoke wait and instead terminates, the child then becomes an **orphan**



---

# Interprocess Communication

---

- ❖ Cooperating - any process that can be affected by another process. Any process that shares data with another process is a cooperating process
- ❖ Independent - any process that cannot affect or be affected by another process



---

# Interprocess Communication

---

- ❖ Why cooperate?
  - ❖ Information Sharing (e.g., shared file)
  - ❖ Computation Speedup (e.g., run in parallel)
  - ❖ Modularity (e.g., system functions in separate processes)
  - ❖ Convenience (e.g., work on multiple tasks at the same time)



---

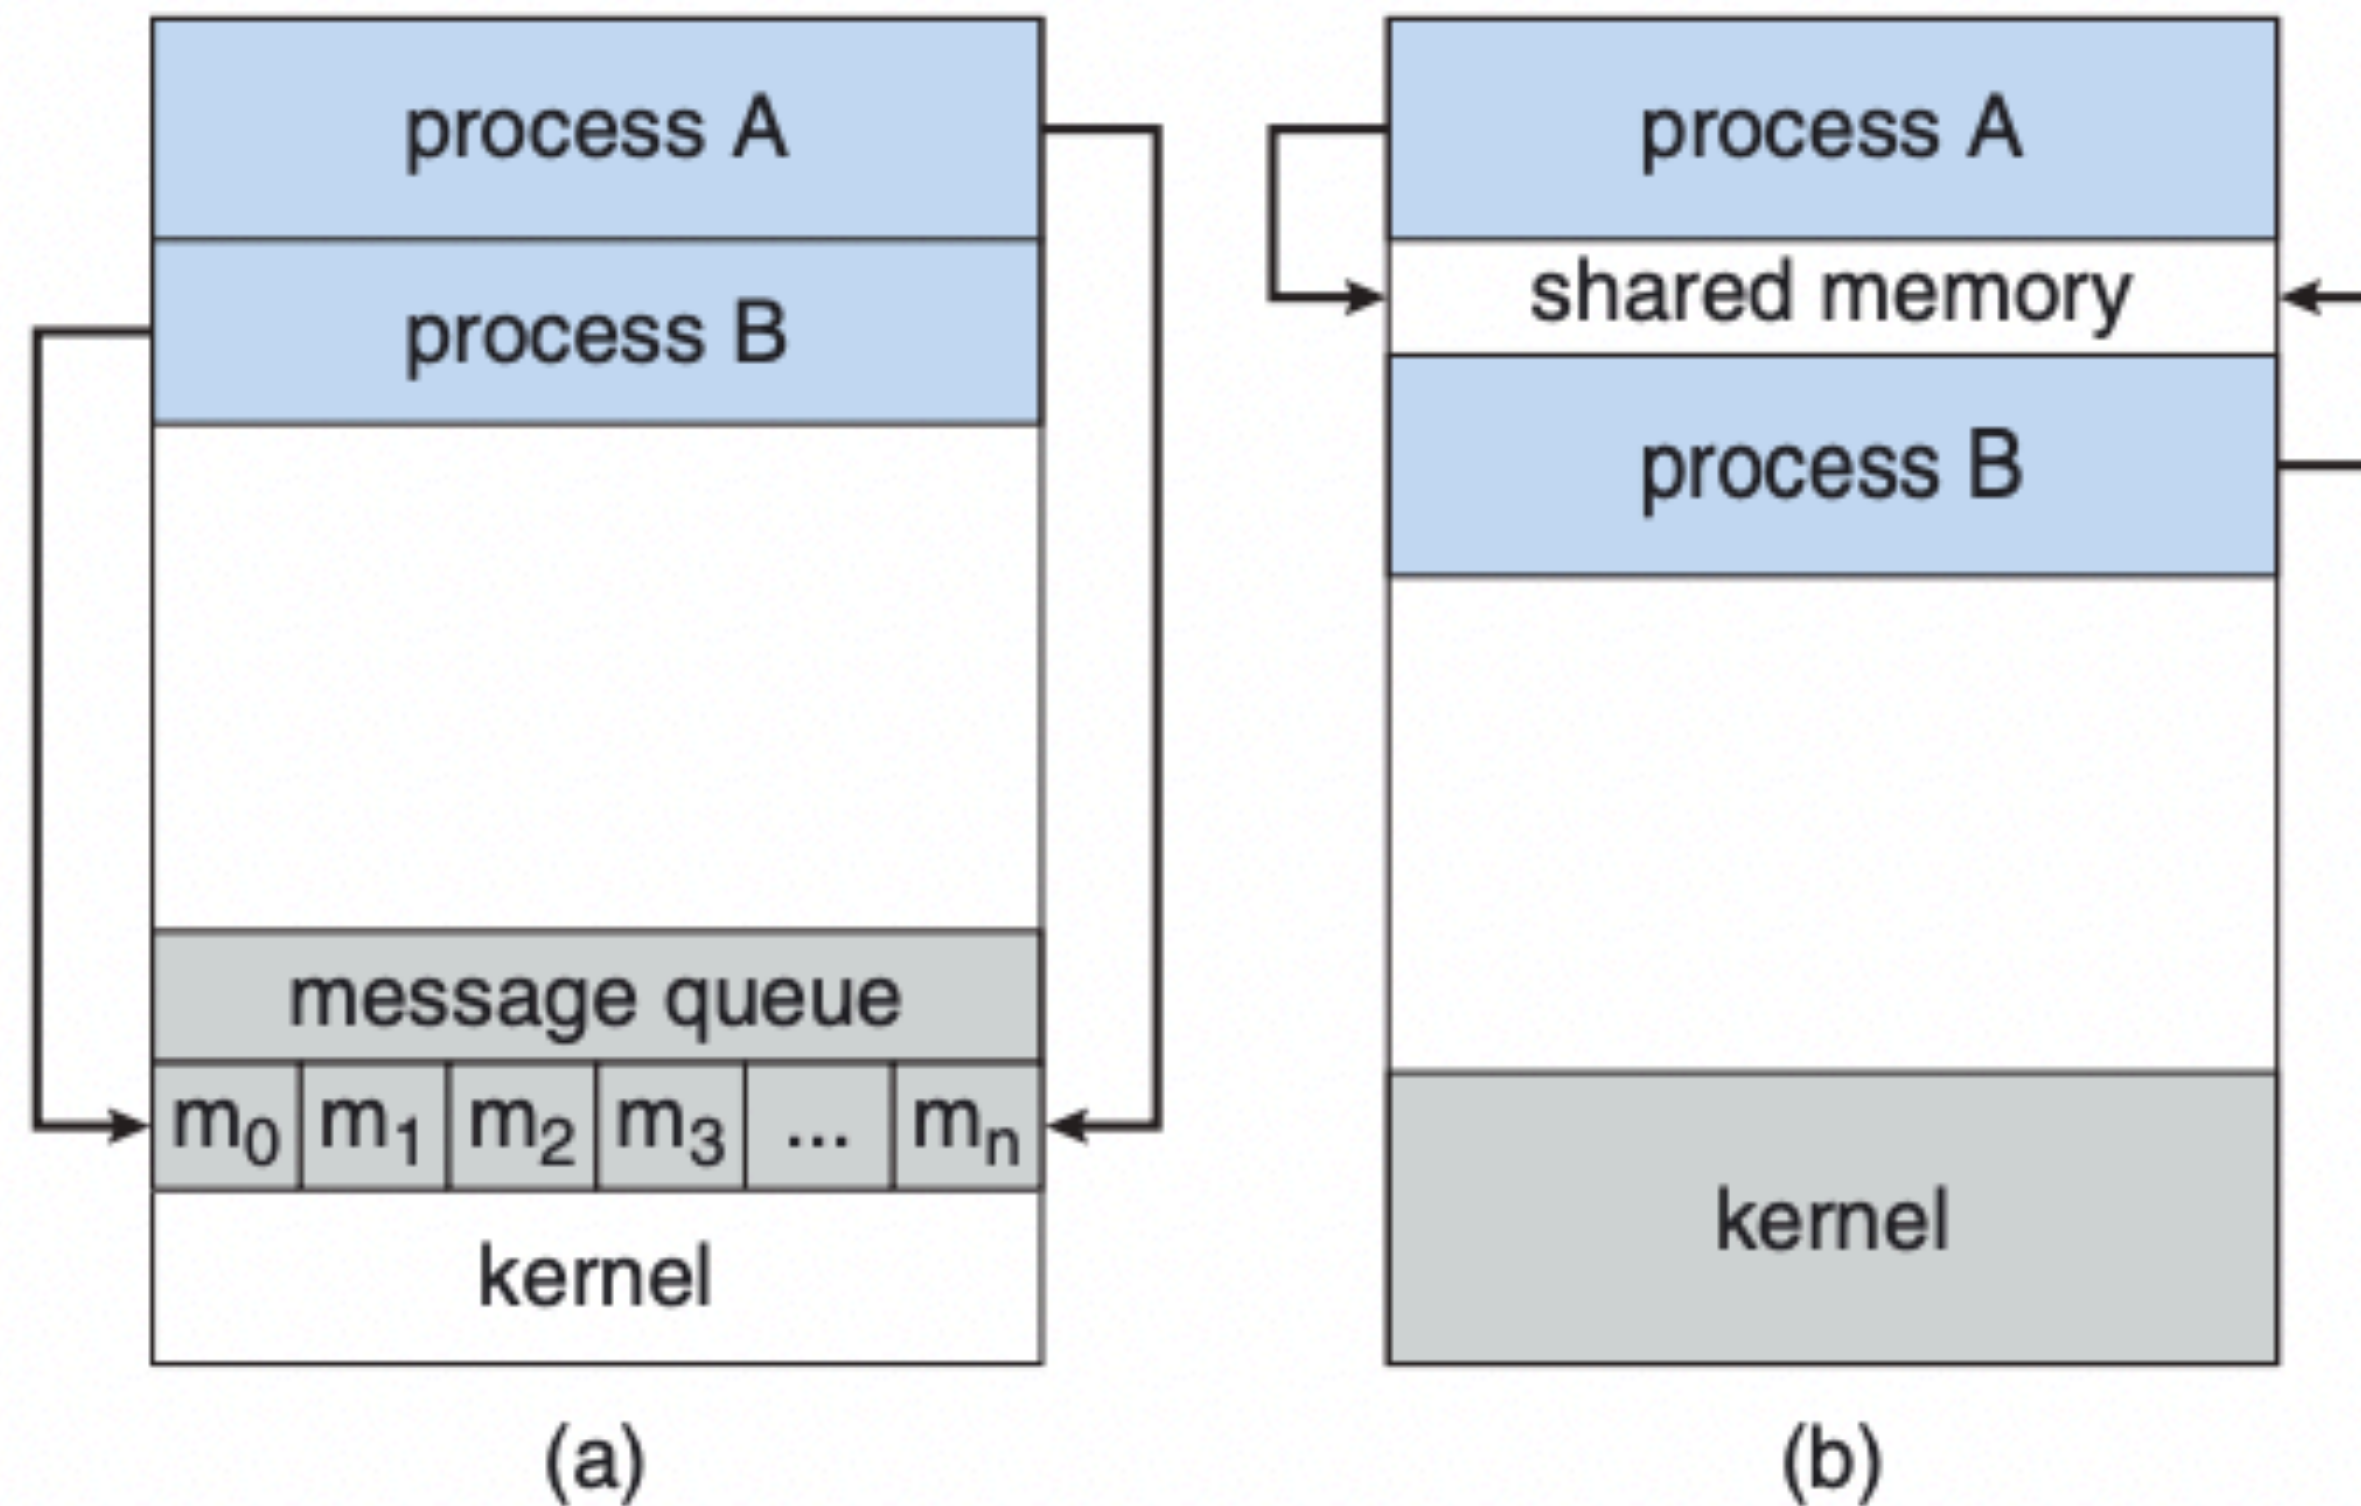
# Interprocess Communication

---

- ❖ Two fundamental models for IPC
  - ❖ Shared Memory
    - ❖ Reading / Writing data to a shared region of memory
  - ❖ Message Passing
    - ❖ Useful for exchanging smaller amounts of data



# Interprocess Communication



**Figure 3.12** Communications models. (a) Message passing. (b) Shared memory.



---

# Shared-Memory Systems

---

- ❖ Create shared memory region
- ❖ Processes attach to the shared address space



---

# Shared-Memory Systems

---

- ❖ Producer / Consumer : One process produces, the other process consumes
  - ❖ What are some examples of a Producer / Consumer scenario



---

# Shared-Memory Systems

---

- ❖ Unbounded buffer - no practical limit on the size of the buffer
  - ❖ Producer can always produce new items
- ❖ Bounded buffer - fixed buffer size
  - ❖ Producer waits if buffer is full
  - ❖ Consumer waits if buffer is empty



---

# Message-Passing Systems

---

- ❖ Two operations:
  - ❖ `send(message)`
  - ❖ `receive(message)`



---

# Message-Passing Systems

---

- ❖ A communication link is established between every pair of processes that want to communicate.
- ❖ A link is associated with exactly two processes



---

# Message-Passing Systems

---

- ❖ Symmetry - sender and receiver name each other to communicate
- ❖ Asymmetry - sender names the recipient but not vice versa
  - ❖ `send(P, message)` - send a message to P
  - ❖ `receive(id, message)` - receive a message from any process, the variable id is set to the name of the process communicated with



---

# Pipe Creation

---

- ❖ Create an array with 2 elements:
  - ❖ `int fd[2]`
- ❖ Create a pipe and return the result of the creation to a variable:
  - ❖ `pipeCreationResult = pipe(fd)`
- ❖ Check to see if the pipe was created successfully, if not, return an error:
  - `if(pipeCreationResult < 0){`
    - `perror("Failed pipe creation\n");`
    - `exit(1);`
  - `}`



---

# Pipe Creation

---

Create a new process using fork, then:

```
if(pid == 0){  
    write(fd[1], &output, sizeof(int));  
    printf("Child wrote [%d]\n", output);  
}  
else{  
    read(fd[0], &input, sizeof(int));  
    printf("Parent received [%d] from child process\n", input);  
}
```



---

# In-Class Assignment

---

Create a pipe as demonstrated in today's lecture.

The program should:

1. Fork off a child
2. Request an input string up to 500 characters
3. Output that the child sent a message (and what it was)
4. Output that the parent received a message (and what it was)